


# ANSY - Getting in Linux Kernel details

version 2022-12-20

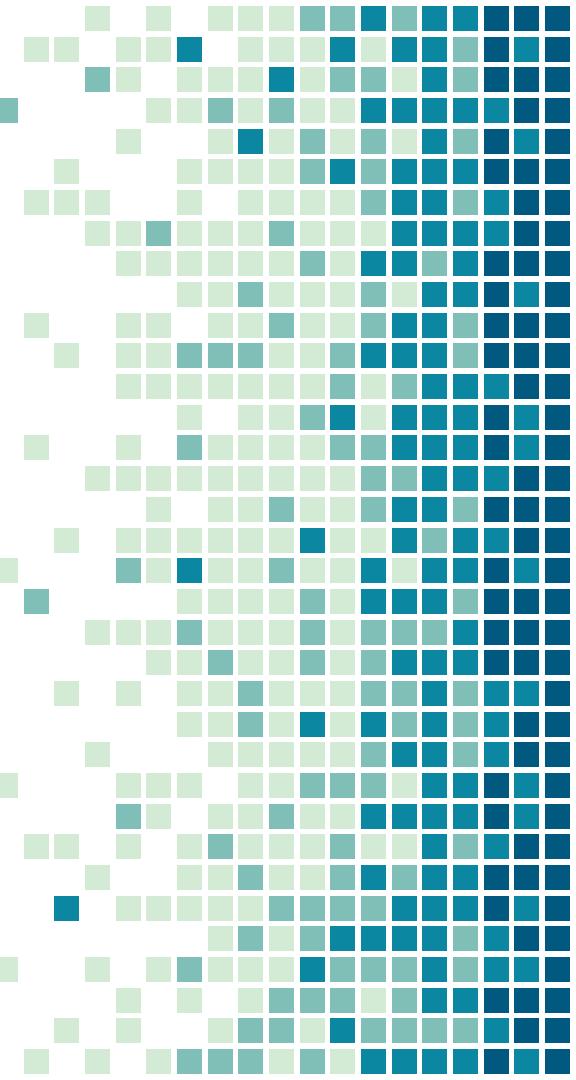


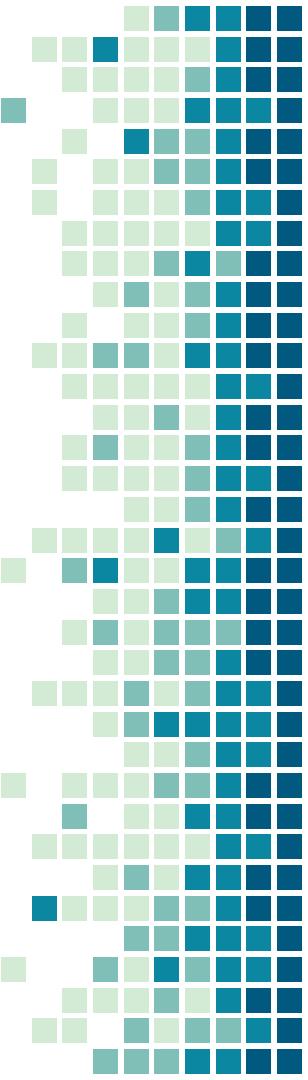
-- Cyril zarak Duval, root CRI/ACU 2020



# Introduction

Generic information about the course





# Why should you listen to the course ?


- Linux is a state of the art in the industry
- Understanding in details will help you as low-level engineers
- Provide some tools useful to work with Linux
- Help you acquire an analytic mind to tackle low-level issues
- Subject somewhat difficult
- Getting a decent grade

# Notions

- ptrace
- strace & rr
- ebpf
- auditd
- kprobes/uprobes/tracepoint
- perf
- pseudo filesystems
- VFS
- iptables/nftables
- initramfs
- PXE
- dracut
- BCC
- CPU & scheduler metrics
- Memory management & metrics
- Systemd
- ...


# Understand observability

Observability applied to Linux




# Observability

- Observability is a high-level notion
- Observability means understanding what is going on in a system
- Observability in its modern approach has 3 pillars:
  - ◆ Metrics
  - ◆ Logs
  - ◆ Tracing
- Linux is the core of all our foundations
- We need to have observability in it
  - ◆ To better understand, to administrate it
  - ◆ To debug



# Observability in Linux

- Linux thankfully offers interfaces for observability
- What are the first things that come to your mind when you think about observability and Linux ?
  - ◆ What are the things you want to observe ?
  - ◆ What are the interface(s) you will use ?



# Get information about CPU usage

Let's start with something "simple"



# CPU and linux


- What does CPU usage means ?
  - ◆ 0, 50, 100% ?
  - ◆ 800% ?
- It a percentage of time spent working on stuff, otherwise idling
- How do we get this number ?
  - ◆ top, htop
  - ◆ mpstat


# CPU and Linux

- How does *they* get the information ?
- Let's pause this question and investigation, and focus on the methodology here

**“** *I have a tool behaving in a way  
that is unknown to me.*

*How do I figure out how it  
works?*





# How does it works ?

- **mpstat** returns CPU usage, along with some useful information
- Does it create this information ?
- Does it collect this information from somewhere ?
  - ◆ Is it on the network ?
  - ◆ Is it on the machine ?
    - Our filesystem ?
    - Any other mean ?

# How do we get this info ?

- 2 hypothesis:
  - ◆ The CPU usage information is returned by the hardware directly
  - ◆ The CPU usage is computed by the kernel and reported
- How can we figure this out ?
  - ◆ Knowing that mpstat *knows* the answer
- The most straightforward solution would be to read mpstat source code
  - ◆ But before actually doing this, let's play a small game

# How do we contact the kernel ?

What are the interfaces offered ?



# Kernel interfaces


- The “only” interface is a syscall
  - ◆ All other high-level interfaces are syscall-based
- A syscall can give us information directly:
  - ◆ `gethostname(2)`
  - ◆ `gettimeofday(2)`
  - ◆ `getcpu(2)`
  - ◆ `getcwd(2)`
  - ◆ ...

# Kernel interfaces

- Some syscalls are used to reach higher level interfaces
  - ◆ open(2), openat(2), read(2), write(2), close(2)
- What are higher level kernel interfaces ?
  - ◆ /proc/....
  - ◆ /sys/....
  - ◆ /dev/...
  - ◆ /sys/kernel/debug/...
  - ◆ /sys/kernel/security/...
  - ◆ /sys/firmware/efi/efivars/...
  - ◆ /sys/fs/cgroup/...

# Kernel interfaces

- Are other syscalls used for higher level interfaces ?
- Yes:
  - ◆ socket(2)
  - ◆ ioctl(2)
  - ◆ bpf(2)
  - ◆ perf\_event\_open(2)
  - ◆ ptrace(2)



# Kernel interfaces

- Let's get back the special directories mentioned before (/sys, /proc, ...)
- How are they special ?
- They aren't "real files" on your SSD
  - ◆ In fact you can open your SSD on another machine and check that by yourself in a very easy and naive way
- The files there are kernel interfaces in the forms of a file
  - ◆ "In UNIX, everything is a file"
- Those are pseudo-filesystems
- But more about that later ...

# Back to our CPU usage analysis

So, what about mpstat ?




# mpstat

- mpstats like almost everything on a classic linux distro is open-source
- Checking source code is therefore a good reflex for things like this
- Let's read [mpstats source code](#)
- Code is well written
  - ◆ Follow many standards
  - ◆ Proper naming convention
  - ◆ Comments
  - ◆ ....

# mpstat

- Code is quite short but finding the information still took time
- Can we make this more efficient ?
- What could be another approach than reading source code ?
  
- What do we know or deduced ?
  - ◆ The information is probably held by the kernel
  - ◆ mpstat gets this information
  - ◆ Communication between userland and kernel land is done via syscalls
  - ◆ Could we just look at the syscalls mpstats did ?



# Let's ptrace mpstat

- Linux offers a syscall and its interface to debug softwares
- ptrace(2)
  - ◆ But more about it later ...
- Used to debug, like GDB, to see what is going on, inspect code, variable values, etc
- What if we have a special debugger ?
  - ◆ This debugger will just run the program
  - ◆ But whenever a function is called, it checks if it a syscall function ?
    - But wait, are syscall functions ?

# Some syscall digression


How is implemented a syscall in the end ?



# What are syscalls ?


- We know that a syscall is a kernel function that is called from userland
  - ◆ Sort of
- But are we allowed to call directly a function like this ?
- In x86 (IA\_32 and x86-64) we run code on the CPU in rings
  - ◆ ring 0 is the most privileged one
    - Allowed to access hardware and configure the CPU directly
  - ◆ ring 3 is the one userland runs in. Can do computation, but cannot run some privileged CPU instructions

# CPU rings



# What are syscalls ?

- The kernel runs in ring 0
- The kernel can therefore do things regular process can't
- Regular process still need to access some protected devices or perform some privileged operations
  - ◆ In a controlled environment (permissions, ....)
- They contact the kernel for those operations via syscalls
- The kernel checks permissions, do sanity checks, etc and performs the operation
- The result, if any, is returned to the user



# What are syscalls ?


- Syscalls are the interface between userland and the kernel
- Allows privileged operation, control kernel behavior or use kernel features
- Allow abstraction (disk drivers, network drivers, ....)
- How can we call a ring 0 function from ring 3 ?
  - ◆ Do we have symbols exported ?
- In x86 we have 3 ways:
  - ◆ INT 0x80 (legacy)
  - ◆ sysenter (IA\_32)
  - ◆ syscall (AMD64)

# Syscall implementation in x86

- INT 0x80 is hardly used anymore, it is a legacy way of making a syscall
  - ◆ Creates an interrupt to notify the kernel
    - But more about that later ...
- sysenter is also called fast system call, created by intel for IA\_32
- syscall is the AMD64 version, mostly used now


# Syscall implementation in x86

- When doing a syscall instruction, what happens exactly ?
- On syscall, the CPU looks in a specific MSR: IA32\_LSTAR
- IA32\_LSTAR MSR contains a ring 0 function address to execute
- In Linux, it's entry\_SYSCALL\_64
  - ◆ [https://elixir.bootlin.com/linux/latest/source/arch/x86/entry/entry\\_64.S#L49](https://elixir.bootlin.com/linux/latest/source/arch/x86/entry/entry_64.S#L49)
- Linux determines which syscall has been called in this function based on %rax
- The arguments to the syscall are in %rdi, %rsi, ... as usual



# Syscall implementation in x86

- In entry\_SYSCALL\_64 we prepare everything to call the actual kernel function
  - ◆ We save userland general purpose registers on the stack and do a few things
- We call the kernel function associated with the syscall requested, forwarding the arguments userland gave
- Once the kernel function returns, we put the return value on the stack
- We do a few things back, put the value from the stack back in registers and return




# Syscall consequences

- Calling a syscall means calling a kernel function through some steps
- The steps are to ensure security when switching from ring 3 code to ring 0 (and vice-versa)
- These steps are called privilege switch
- Privilege switch is quite expensive:
  - ◆ More instructions to execute
  - ◆ No CPU pipelining/branch prediction/...
  - ◆ Data shouldn't be accessed directly and shall always go through `copy_from_user`/`copy_to_user`
  - ◆ Pointers must be handled carefully

# About vDSO

Some more information about syscall implementation




# Virtual Dynamic Shared Object

- Some syscalls are used a lot
- They don't have any security and will behave the same for every user (privileged or not)
- To allow better performances, the kernel exposes some syscall directly in userland
- Userland implementation is done in vDSO
- Shared ELF object to every userland process
  - ◆ Address fetchable via auxiliary values
  - ◆ Glibc gets it for you

# Virtual Dynamic Shared Object

- Contains some syscall, depending on the architecture
- gettimeofday(2), getcpu(2), ...
- No privilege switch = faster
  - ◆ No strace, no seccomp however




# Virtual Dynamic Shared Object

```
1 $ file vdso64.so
2 vdso64.so: ELF 64-bit LSB shared object, x86-64, version 1 (SYSV), dynamically linked,
  BuildID[sha1]=1e5d0c12baeaaebcc0ca26b4c4c4dd0be8f0e5d4, stripped
3 $ objdump -T vdso64.so
4
5 vdso64.so:      file format elf64-x86-64
6
7 DYNAMIC SYMBOL TABLE:
8 0000000000000940  w  DF .text  00000000000000264  LINUX_2.6    clock_gettime
9 0000000000000760  g  DF .text  0000000000000001b0  LINUX_2.6    __vdso_gettimeofday
10 0000000000000bb0  w  DF .text  000000000000000063  LINUX_2.6   clock_getres
11 0000000000000bb0  g  DF .text  000000000000000063  LINUX_2.6   __vdso_clock_getres
12 0000000000000760  w  DF .text  0000000000000001b0  LINUX_2.6  gettimeofday
13 0000000000000910  g  DF .text  00000000000000002e  LINUX_2.6   __vdso_time
14 0000000000000c50  g  DF .text  0000000000000000a1  LINUX_2.6   __vdso_sgx_enter_enclave
15 0000000000000910  w  DF .text  00000000000000002e  LINUX_2.6   time
16 0000000000000940  g  DF .text  00000000000000264  LINUX_2.6   __vdso_clock_gettime
17 0000000000000000  g  DO *ABS*  00000000000000000000  LINUX_2.6   LINUX_2.6
18 0000000000000c20  g  DF .text  00000000000000002a  LINUX_2.6   __vdso_getcpu
19 0000000000000c20  w  DF .text  00000000000000002a  LINUX_2.6   getcpu
```

# Let's get back to mpstats system calls

Now we do know what a system call is



# Let's ptrace mpstat

- Linux offers a syscall and its interface to debug softwares
- ptrace(2)
- Used to debug, like GDB, to see what is going on, inspect code, variable values, etc
- What if we have a special debugger ?
  - ◆ This debugger will just run the program
  - ◆ But whenever a function is called, it checks if it a syscall function
  - ◆ If so, prints arguments, resume execution and print return value

# Discovering strace

- A famous debugging tool for such purposes exists
- strace(1)
- In the simplest usage:
  - ◆ Starts a process with given arguments
  - ◆ Gets notified of all the syscalls the tracee performs
  - ◆ Prints the syscall, its arguments and return value


# Usefulness of strace


- When is strace useful ?
  - ◆ Find out why a software fails
  - ◆ Find out how it behaves if no documentation
    - Ex: location of config files read by the app
    - Ex: Interaction with other processes
    - Ex: Memory impact and behaviour
  - ◆ See where a software hangs (if on a syscall)
  - ◆ ....
- strace is a very popular and versatile debug tool
  - ◆ Simpler and quicker to use than GDB
    - Not suited for all workflows though

# How to strace ?

- How to use strace(1) efficiently ?
- A few tips:
  - ◆ Use -f to follow and strace forks too
  - ◆ Use -z or -Z to see only successful or failed syscalls
  - ◆ -c will give you a summary/overview of the syscalls used. Can be useful at first to have a sneak peak
  - ◆ Discover the -e option

“ Let's discover strace -e and play  
with strace a bit







# Strace can be difficult

- Since strace will show every syscalls, it might be difficult to find what you're looking for
- Especially if the software is huge
  - ◆ Or we have limited knowledge on what to look for
- Example: pylint
  - ◆ Where's the configuration file ?
  - ◆ Tracing open(2)/openat(2) ?

# Strace


5 % time	seconds	usecs/call	calls	errors	syscall
7 45.49	0.006295	0	19086	9931	newfstatat
8 12.04	0.001666	0	1677		read
9 11.13	0.001540	12	128		munmap
10 9.31	0.001288	1	991	10	openat
11 4.94	0.000683	1	433		brk
12 3.91	0.000541	0	985		close
13 3.66	0.000506	0	1534	3	lseek
14 2.51	0.000347	1	273		mmap
15 2.29	0.000317	6	48		write
16 2.19	0.000303	0	899	891	ioctl
17 1.65	0.000228	2	106		getdents64
18 0.30	0.000041	1	32		mprotect
19 0.26	0.000036	0	66		rt_sigaction
20 0.12	0.000017	1	16	9	readlink
21 0.09	0.000012	0	53		getcwd
22 0.05	0.000007	1	5		getrandom
23 0.03	0.000004	4	1		sysinfo
24 0.02	0.000003	3	1		futex
25 0.02	0.000003	1	2		prlimit64
26 0.01	0.000002	0	3		dup
27 0.00	0.000000	0	2		pread64
28 0.00	0.000000	0	1	1	access
29 0.00	0.000000	0	1		execve
30 0.00	0.000000	0	1		fcntl
31 0.00	0.000000	0	1	1	mkdir
32 0.00	0.000000	0	2	1	arch_prctl
33 0.00	0.000000	0	1		gettid
34 0.00	0.000000	0	1		set_tid_address
35 0.00	0.000000	0	1		set_robust_list
36 0.00	0.000000	0	1		epoll_create1
37 0.00	0.000000	0	1		rseq
38					
39 100.00	0.013839	0	26352	10847	total





# Strace can be difficult

- Example: pylint
  - ◆ Where's the configuration file ?
  - ◆ Tracing open(2)/openat(2) ?
    - 991 openat(2) in this example
  - ◆ Maybe grepping "cfg", "yml" or "json" ?
  - ◆ Actually file name is pylintrc
  - ◆ And not even open/openat(2) if doesn't exist !



# Strace can be difficult



```
1 $ strace pylint ipxe_manager 2>&1 | grep pylintrc
                                         30 ↵
2 newfstatat(AT_FDCWD, "pylintrc", 0x7ffcd8fc3b10, 0) = -1 ENOENT (No such file or directory)
3 newfstatat(AT_FDCWD, ".pylintrc", 0x7ffcd8fc3b10, 0) = -1 ENOENT (No such file or directory)
4 newfstatat(AT_FDCWD, "/home/zarak/.pylintrc", 0x7ffcd8fc3b10, 0) = -1 ENOENT (No such file or directory)
5 newfstatat(AT_FDCWD, "/home/zarak/.config/pylintrc", 0x7ffcd8fc3b10, 0) = -1 ENOENT (No such file or
   directory)
6 newfstatat(AT_FDCWD, "/etc/pylintrc", 0x7ffcd8fc3cd0, 0) = -1 ENOENT (No such file or directory)
7 newfstatat(AT_FDCWD, "pylintrc", 0x7ffcd8fc6c70, 0) = -1 ENOENT (No such file or directory)
8 newfstatat(AT_FDCWD, ".pylintrc", 0x7ffcd8fc6c70, 0) = -1 ENOENT (No such file or directory)
9 newfstatat(AT_FDCWD, "/home/zarak/.pylintrc", 0x7ffcd8fc6c70, 0) = -1 ENOENT (No such file or directory)
10 newfstatat(AT_FDCWD, "/home/zarak/.config/pylintrc", 0x7ffcd8fc6c70, 0) = -1 ENOENT (No such file or
    directory)
11 newfstatat(AT_FDCWD, "/etc/pylintrc", 0x7ffcd8fc6e30, 0) = -1 ENOENT (No such file or directory)
```



# Strace can be difficult

- Useful sometimes to simulate a failure or to simulate a success of a syscall or a set of syscalls
- Need to have some knowledge of the software
  - ◆ Or some intuition

# CPU metrics

What is a CPU with linux ?



# Understanding CPU metrics

- A CPU core or thread shall already be known to you
- We've seen already 2 kinds of things a CPU can execute in this course:
  - ◆ User code
  - ◆ Kernel code
- What are the other things a CPU can do ?
- Fortunately a CPU isn't always doing something: it can idle
- Let's check the metrics exported by the kernel in /proc/stat

# Understanding CPU metrics


- User
- Nice
- System
- Idle
- Iowait
- Irq
- Softirq
- Steal
- Guest
- Guest\_nice

# Understanding CPU metrics

- User -> userland code
- Nice
- System -> kernel-land code
- Idle -> CPU literally doing nothing (~no power usage, C-state)
- Iowait
- IRQ
- SoftIRQ
- Steal
- Guest
- Guest\_nice

# Understanding CPU metrics

- User -> userland code
- Nice
- System -> kernel-land code
- Idle -> CPU literally doing nothing (~no power usage, C-state)
- Iowait
- IRQ
- SoftIRQ
- Steal
- Guest -> kernel KVM gave CPU time to VM
- Guest\_nice -> kernel KVM gave nice CPU time to VM




# Process niceness and scheduler

Why isn't pulseaudio nice ?




# CPU and multithreading

- A classic PC/server runs dozens if not hundreds of processes in “parallel”
- A modern CPU has multiple cores, and multiples threads or logical cores/hyper-thread
- Let's say our CPU has 16 logical cores
- I can truly execute 16 processes in parallel
- How can I give the impression it's running 150 ?




# CPU and multithreading

- Most processes don't need the CPU 100% of the time
- They need some time to work, and have to wait
  - ◆ Timer, user input, IO, being activated back, ...
- If most of them don't need to actually run in parallel, we can split execution in small timeshares, and simulate parallel execution
- This is the role of the scheduler to provide such timeshares and execute processes




# CPU and multithreading

- Here is the classical representation of a Process state in Linux




# Process states

Giving meaning to R/S/D/Z/T



# CPU and multithreading

→ Actually it looks more like this



# Process state

- R state means running or runnable
  - ◆ Either currently being executed on a CPU core (running)
  - ◆ Or waiting for a core to be free and for the scheduler to start it (runnable)
- S state is the state some process will spend the most time in
  - ◆ Waiting for an event, for I/O, for a timer, ...
- T state is fairly easy to grasp, one stopped the process by sending a SIGSTOP signal

# Process state


- D state is a bit more shady
  - ◆ Some linux syscall are not interruptible. It means that a process waiting for the syscall to complete cannot be killed.
  - ◆ No signal can be transmitted, even SIGKILL
  - ◆ Examples include some I/O syscalls, KVM related calls, etc
    - [https://elixir.bootlin.com/linux/latest/A/ident/TASK\\_UNINTERRUPTIBLE](https://elixir.bootlin.com/linux/latest/A/ident/TASK_UNINTERRUPTIBLE)
  - ◆ Famous example often found is a NFS-related process stuck in D-state when NFS server is unreachable

# Process state

- Z state is for a zombie
  - ◆ Zombie process is a process that has finished its execution but hasn't been wait(2)-ed by its parent
  - ◆ Its information remains and must be collected for the process to be removed from the process list
  - ◆ Init process must wait for zombie process re-attached to it to maintain a clean system
  - ◆ When a zombie is created, SIGCHLD is sent to parent process

# What does it have to do with niceness ?

Exploring CFS




# CFS - Completely Fair Scheduler

- The role of the process scheduler is to run process when it makes sense
  - ◆ When they are ready to run
  - ◆ When they can (i.e. a CPU core is available)
- Linux default scheduler is called CFS
- It divides time in timeslices
- It gives a timeslice to the process that is ready to be run and has been starving CPU time the most first

# CFS - Completely Fair Scheduler


- If the system is not overloaded, CFS doesn't have to make important decisions
  - ◆ Most process are in D/S state, and therefore very few in R state. Decisions are easy
- But if the system starts to be overloaded, CFS comes to play
- CFS selects the process that is missing the most vruntime, i.e. the process that should have been running but hasn't
  - ◆ Takes decision based on total execution time and how long it has been waiting

# CFS - Completely Fair Scheduler



- CFS tracks process via a red/black tree
  - ◆ On the left of the tree, process with the smallest vruntime
- It is also able to dynamically change the length of the CPU timeslice based on the load:
  - ◆ If a process is alone, it makes sense to give it a lengthy timeslice since it won't impact anyone
  - ◆ If 2 process requires each 50% of a single CPU core, to make them look like they run in parallel we need to alternate their execution
    - But mind context\_switch ! Intervenes sched\_min\_granularity\_ns

CFS



```
struct task_struct {
    volatile long state;
    void *stack;
    unsigned int flags;
    int prio, static_prio normal_prio;
    const struct sched_class *sched_class;
    struct sched_entity se; ←
    ...
};
```


```
struct ofs_rq {
    ...
    struct rb_root tasks_timeline;
    ...
};
```

```
struct sched_entity {
    struct load_weight load;
    struct rb_node run_node;
    struct list_head group_node;
    ...
};
```

```
struct rb_node {
    unsigned long rb_parent_color;
    struct rb_node *rb_right;
    struct rb_node *rb_left;
};
```

# CFS - Completely Fair Scheduler

- But CFS is more complex than that
- Some process needs higher priority in their scheduling, because scheduling latency impacts
  - ◆ I.e. audio
    - audio doesn't need a lot of CPU time
    - But audio suffers heavily from latency
- This is the niceness of a process with linux
- The nicer the process, the less priority it gets
- Very nice process can still take 100% of a CPU core. They will just be descheduled if anyone else is asking for some CPU time



# CFS - Completely Fair Scheduler

- CFS also takes into account various other configuration
- Internally, it also has a concept of priority
- Priority is changed by niceness, but to a range only
- To access the other priority values, a process must change its scheduling class
- More info on sched(7)

# CPU load

A metric often misunderstood



# CPU load

- CPU load can be understood as "how many operations my CPU is currently doing"
- This is a wrong understanding when it comes about the load metric reported by linux
- A better understanding would be "How much pressure is being applied to the CPU in average for a period of time"
- What does it measure exactly ?

# CPU load

- Linux load represents the number of processes running, or waiting to be ran on the system, in average for a period of time
  - ◆ It also includes processes in uninterruptible sleep
    - I/O matters
  - ◆ It is not limited to a core -> all load values don't have the same meaning on each machine
  - ◆ Usually troubles begins when the load reaches the number of CPU cores

# CPU load

- 3 values exported in /proc/loadavg
  - ◆ 1 min, 5min and 15min load
  - ◆ Number of processes in R state / schedulable entities
  - ◆ PID of the latest created process



```
1 $ cat /proc/loadavg
2 1.34 1.68 1.63 3/3253 2817976
```

# CPU load


- Having 3 loads metrics, and them being averages has impact
- There is delay between event and possible visualization on the curves
- load1 is closer to “instant” load while load15 is really difficult to pull in any direction

# CPU load

Load over time


Start of event

End of event




# Let's get back to our CPU metrics

We do know now who's nice and who  
isn't




# Understanding CPU metrics



- User -> userland code
- Nice -> process with high niceness
- System -> kernel-land code
- Idle -> CPU literally doing nothing (~no power usage, C-state)
- Iowait
- IRQ
- SoftIRQ
- Steal
- Guest -> kernel KVM gave CPU time to VM
- Guest\_nice -> kernel KVM gave nice CPU time to VM

# Understanding CPU metrics

- User -> userland code
- Nice -> process with high niceness
- System -> kernel-land code
- Idle -> CPU literally doing nothing (~no power usage, C-state)
- Iowait
- IRQ
- SoftIRQ
- Steal -> As a VM, hypervisor didn't schedule us
- Guest -> kernel KVM gave CPU time to VM
- Guest\_nice -> kernel KVM gave nice CPU time to VM



# Understanding CPU metrics

- User -> userland code
- Nice -> process with high niceness
- System -> kernel-land code
- Idle -> CPU literally doing nothing (~no power usage, C-state)
- Iowait -> time spent for a process waiting for I/O (unreliable)
- IRQ
- SoftIRQ
- Steal -> As a VM, hypervisor didn't schedule us
- Guest -> kernel KVM gave CPU time to VM
- Guest\_nice -> kernel KVM gave nice CPU time to VM

# Quick tour of IRQ and softIRQ

Bringing some memory back



# CPU interrupts

- The way for the hardware to notify the CPU something is happening is through IRQ
- For example, the user moved its mouse or typed on its keyboard
- Paquets reached the machine and are waiting on the network card
- Without getting into too many details, the CPU gets notified of these events through the PIC (Programmable Interrupt Controller)

# CPU interrupts (x86)

- An interrupt stops the current CPU execution and executes an interrupt handler read on the IDT
  - ◆ The IDT (Interrupt Descriptor Table) maps interrupts to handlers
- An interrupt can be triggered by external device (like the network card) or by the CPU itself
  - ◆ In this case it's called a software interruption
    - Or an Exception (x86)
  - ◆ Examples include a division by 0, or an INT instruction

# CPU interrupts (x86)

- Exceptions (or software interrupts) are of 3 categories:  
Traps, Fault and Abort
  - ◆ A trap is reported after the execution (ex: INT) and allow process continuity
  - ◆ A Fault is reported before the actual execution to allow to fix it (ex: div / 0)
  - ◆ An Abort is when everything is on fire. Run.
  - ◆ More about it in the x86 Intel manual

# CPU interrupts (x86)

- In linux, it's translated as interrupts/IRQ (Interruption ReQuest) and softIRQ (software IRQ)
- Values are exposed in /proc/interrupts
- softIRQ in linux don't show all x86 exceptions
  - ◆ softIRQ displayed by Linux are limited, check /proc/interrupts
  - ◆ softIRQ is a "primitive" system that has been partially taken over by tasklets
- There is no direct mapping between linux exposed values and x86 events

# CPU metrics in the end


Putting everything together



# Understanding CPU metrics

- User -> userland code
- Nice -> process with high niceness
- System -> kernel-land code
- Idle -> CPU literally doing nothing (~no power usage, C-state)
- Iowait -> time spent for a process waiting for I/O (unreliable)
- Irq -> hardware interrupts
- Softirq -> software interrupts
- Steal -> As a VM, hypervisor didn't schedule us
- Guest -> kernel KVM gave CPU time to VM
- Guest\_nice -> kernel KVM gave nice CPU time to VM


# Understanding CPU metrics



- htop CPU bars have colors representing the different kind of CPU metrics
- By default:
  - ◆ blue = nice
  - ◆ green = user
  - ◆ red = kernel (+ iowait + irq + softirq)
  - ◆ orange = guest (+ steal)

# PSI - how to represent pressure

Getting a higher level metric to abstract





# Monitor system going wrong

- Let's say you want to monitor your system and try to detect problematic states
- What is a problematic state ?
  - ◆ Let's define this in this context by "a state when your workload doesn't run properly or in a degraded state, not exploiting your machine full capacity"
- In this case, is a 100% CPU usage defined as a problematic state ?

# Understanding when a state is problematic

- CPU is a complex metric to grasp when trying to investigate problematic situations
  - ◆ Quite some metrics
  - ◆ Easy to get fooled
    - “OMG my CPU is spending all its resources on idle !!”
- CPU used at 100% doesn't mean your application is disturbed
- In some cases, it can be impacted without reaching 100%
- Try to put your metrics in correlation with your application





# Understanding when a state is problematic

- 100% CPU usage when compiling the kernel
  - ◆ Usually not a problem
    - Indicates reaching your max capacity. You might want to upgrade your CPU maybe ?
  - ◆ Can be if done alongside other workload
  - ◆ Niceness to keep in mind
- It's better to rely on what you observe
  - ◆ Latency, mouse lag, etc
  - ◆ How to program this ?

# Monitor system going wrong

- Load is an indicator indeed but:
  - ◆ Relative value (number of cores)
  - ◆ No indication of actual waiting time the process had to wait
  - ◆ R+D state, so few faulty NFS process and the load goes up the roof
  - ◆ Average over time

# Monitor system going wrong

- Linux proposes another metric: Pressure Stall Information
- From ~2018 by Facebook
- 3 metrics: io, memory and CPU
- Represent the % of time wasted because of processes conflicts for a resource
  - ◆ You can have 100% used CPU core and 0% CPU PSI
- Has a polling interface
  - ◆ Used to loadbalance workload in Facebook


# What about memory ?

Yet another complex metric



# Different kinds of memory

- Memory is a wide term with different kinds:
  - ◆ Volatile, fast memory (RAM)
  - ◆ Non-volatile, slower memory (swap)
- When trying to understand memory for your system, 2 kinds:
  - ◆ Virtual memory
  - ◆ Physical memory
- The kernel in combination with the MMU (Memory Management Unit) is responsible for abstracting memory to userland




# Physical memory

- Physical memory is divided in multiple places
  - ◆ I.e. 4x16 GiB of RAM in 4 sticks
  - ◆ 4 GiB of swap on your NVME disk
- Physical memory has its own address space
  - ◆ Depends on the lanes you're plugging the memory in, the motherboard, ...
- Different sources of memory may have different latencies
- ....

# Virtual memory

- Userland doesn't want to deal with this
- Userland wants a unique address space for memory
- For security reasons, userland processes must not be able to access memory from each other
  - ◆ Per process address space
- Userland wants the kernel to do things for it
  - ◆ Maybe he wants to interfere a bit with the decisions
    - Advise, flush, ...
    - Control memory-related mechanisms (i.e. swap)



# Virtual memory

- Memory will be used in multiple cases:
  - ◆ Variable storage (or generic computation needs)
  - ◆ Process executable binary
  - ◆ Disk cache
  - ◆ Kernel memory
  - ◆ Page tables
- Memory is used with pages
  - ◆ 4 kiB on AMD64
  - ◆ Possibility to increase with THP (Transparent Huge Pages)

# Process virtual memory

How do a process have access to  
“memory”?



# Process virtual memory

- Each process has its own address space
- For obvious security reasons
- Each process address space is virtual
  - ◆ 2 process can share the same address in their virtual memory that leads to completely different “real” memory
- Each process address space is flat: no segmentation
- Different sections in their address space however
  - ◆ Everything is not identical

# Process virtual memory

- Example include a process' own executable code in a memory map, called text
- A process global variables in a data section
- ....
- Each process has a struct mm\_struct to describe their virtual address space
  - ◆ Actually threads share the same struct since they have the same address space
  - ◆ [https://elixir.bootlin.com/linux/v6.0.7/source/include/linux/mm\\_types.h#L486](https://elixir.bootlin.com/linux/v6.0.7/source/include/linux/mm_types.h#L486)

# Process virtual memory

- Various interesting implementation details about struct mm\_struct
  - ◆ Like mmap & mm\_rb fields
- Each process can have (and actually have) VMAs
- Virtual Memory Area
- Implements an area of virtual memory, with its property
- struct vm\_area\_struct
- [https://elixir.bootlin.com/linux/v6.0.7/source/include/linux/mm\\_types.h#L397](https://elixir.bootlin.com/linux/v6.0.7/source/include/linux/mm_types.h#L397)

# Process virtual memory

- A VMA is associated to a mm\_struct
- It has flags (including R / W / X)
- Makes the link to a file (if not anonymous memory)
- VMAs can be seen in /proc/<pid>/maps
- Each call to mmap creates a VMA
  - ◆ Sort of, sometimes if it possible, there are merged together

# Pages and Huge pages

Lots of pages and yet not a book released



# Virtual memory – pages

- A page of 4 kiB means 256 000 pages for 1 GiB of RAM
  - ◆ Memory overhead
  - ◆ TLB – Translation Lookaside Buffer
- Possibility to have bigger pages to reduce costs
  - ◆ 2 MiB instead of 4 kiB -> 512 times less TLB entries
- Can be “dangerous”
  - ◆ Pages allocated but not used don’t count
  - ◆ You can malloc(1024 \* 1024 \* 1024); no RAM will be taken
  - ◆ You write 1 bit, the page is actually taken
  - ◆ You will waste more with huge pages

# Virtual memory - THP

- THP can be disabled system wide
  - ◆ Or only used when explicitly asked with madvise(2)
- madvise(2) indicates what usage a portion of memory will be subjected to
  - ◆ The kernel will apply optimization for such usage
- Usage includes:
  - ◆ Normal
  - ◆ Random
  - ◆ Sequential
  - ◆ willneed/dontneed
  - ◆ (un)mergeable (KSM)
  - ◆ (no)hugepage

# Virtual memory - THP

- People usually don't care nor pay attention to THP
- Many applications will malloc( $4 * 1024$ ) 512 times instead of allocating ( $4 * 1024 * 512$ ) directly
  - ◆ Most never use madvise either
- Therefore, THP system would be unused and useless
- Linux introduces khugepaged and heuristics
  - ◆ When allocation more than 2 MiB of RAM at once, you usually allocate a THP without knowing it
  - ◆ khugepaged will look for pages to merge in a THP

# Let's talk about swap

Why swap is bad but still good (??)



# Swap memory

- For most people, memory = RAM
  - ◆ It's fast, usually big enough
  - ◆ However volatile, so we need to be careful
- However RAM is pricey
  - ◆ It's quite easy to run out of RAM even with normal (but relatively heavy) applications/processes
- What happens when we run out of RAM ?
  - ◆ If it's the only memory: oom-killer
  - ◆ If not: swap (then oom-killer if we abuse)

# Swap memory


- Use a persistent disk as a backing storage for more memory
- Disk can be of various nature (various costs and speed)
  - ◆ HDD
  - ◆ SSD
  - ◆ NVME
  - ◆ the shitty 1 GiB USB 1.0 key Capgemini or Sopra Steria gave you in exchange for a CV no one will take time to read
- Performances will suffer: swap is used as a last resort option

# Swap memory

- When do we swap ?
- Swap is used when memory pressure is high. You will not use swap before reaching a huge RAM usage first
  - ◆ Swapped memory stays in swap if unaccessed even if the system memory goes down again
- What are the consequences ?
  - ◆ Swapping out process that takes CPU time and I/O
  - ◆ "Volatile memory" written to non-volatile device (!)
  - ◆ (very) slow memory access on swapped out memory

# Swap memory

- Where on disk is stored swap ?
- 2 options:
  - ◆ Dedicated partition on a disk
  - ◆ Dedicated file on your filesystem
    - Must be on a persistent storage (no tmpfs, duh)
- Swap device needs a specific partition type (`mkswap(8)`)
- Can be enabled and disabled on runtime with  
`swapon(2)(8)/swapoff(2)(8)`



# Swap memory

- Why is swap bad ?
  - ◆ Once you start swapping, performances goes down the drain (-ish)
  - ◆ .. ?
- Why swap is needed ?
  - ◆ Swap isn't used if the system isn't stressed
  - ◆ Most OS or applications don't have an efficient way to react on memory pressure to free-up memory
  - ◆ Reaching the limits often means bad things
  - ◆ There are traps when reaching high memory usage

# Swap memory – vm.swappiness

- What pages of memory shall be sent to the disk ?
  - ◆ What are the best candidates ?
- Intuitively one will say:
  - ◆ Memory rarely accessed
  - ◆ Memory rarely written
  - ◆ LRU
- On top of those cases, one important case to not miss is the memory file-backed, or non-anonymous memory
  - ◆ i.e.: `mmap()` of a file, a process binary, ...
  - ◆ This memory is recoverable. We can evict it safely from memory altogether

# Swap memory - vm.swappiness

- Behavior of what to do when running out of performing memory is controllable via `vm.swappiness`
- The value range goes from 0 to 200 (recent kernel)
- It is often misunderstood

# Swap memory – vm.swappiness

- What people think (it's wrong):
  - ◆ vm.swappiness goes from 0 to 100
  - ◆ It indicates the memory threshold at which the kernel will start swapping
  - ◆ I.e: if vm.swappiness is at 60, if you take 59% of RAM, no swap, 61%, it will start swapping (maybe)
- This is stupid and wrong
  - ◆ Why 60 % ? Why would the kernel voluntarily drop performances to swap ?

# Swap memory - vm.swappiness


- What it does:
  - ◆ It's a balance pressure indicator to put more pressure on swapping out anon pages and dropping recoverable file pages
  - ◆ It's from 0 to 200. 0 means aggressive on file pages, 200 on anon pages
  - ◆ The pressure finally applied is a bit complicated:
    - Swappiness is ignored in some cases
    - Pressure balance is ignored for some part of the list to ensure no leftovers
    - Swappiness = 0 -> no swap unless big troubles

# Swap memory - vm.swappiness

- More information on  
<https://elixir.bootlin.com/linux/v6.0.6/source/mm/vmscan.c#L2731>

# Fear the OOM-killer

How linux kills userland processes *by design*



# OOM-killer

- When reaching the final limit of available memory
- Kernel mechanism triggered on allocation failure
- Find the most suited process to kill
  - ◆ Highest oom\_score
- What is oom\_score ?
  - ◆ Per process score always maintained
  - ◆ Amount of RAM being taken
  - ◆ oom\_score\_adj
  - ◆ Used to be more complex (user vs root process, HW direct access, ...)

# OOM-killer


- Configurable in /proc (like most kernel mechanisms)
  - ◆ Can be disabled entirely
- Killing processes by design “omg wtf”
  - ◆ What do you expect from a system running out of memory anyway ?
- Invocation and its actions logged in /dev/kmsg (dmesg(1))

# Getting rid of the OOM-killer

- Some people don't like the OOM-killer
  - ◆ But still reckon the job is useful
- Namely for a major reason: it intervenes when it's late
  - ◆ Often too late
- 3 projects exists to basically do the same thing, but in userland:
  - ◆ earlyoom
  - ◆ lmkd on Android
  - ◆ Systemd-oomd (which uses memory PSI)

# Kernel threads

You're not the only one having daemons



# Kernel threads

- The kernel does some tasks synchronously:
  - ◆ Syscalls kernel code is executed when the user calls it
- But there are also asynchronous tasks to perform:
  - ◆ Kswapd for example will swap out memory
    - Even compress it with zcache enabled
  - ◆ khugepaged does periodic scans to reduce memory fragmentation by merging pages in THP
- Kernel threads are visible with ps or htop like other processes
- They don't have an associated mm\_struct

# Understanding memory metrics

Why is my process taking 17 GiB of RAM  
on my 16 GiB laptop ?




# Memory metrics

- The memory metrics we're the most interested in for basic usage is free memory
  - ◆ In fact it's incorrect. We want to know what memory is still usable for our process
- Free memory != Usable memory
  - ◆ Caches
  - ◆ Buffers
- When checking for available memory with `free -m` for example, be careful to read "available" and not "free"

# Memory metrics

- The interface to check global memory usage for the machine is /proc/meminfo
- It lists memory and breaks it in different kind of usage
- It's has a lot of fields, some of them are overlapping or imprecise
- It can be misleading and quite difficult to understand it




# Memory metrics /proc/meminfo

→ Example on my 24 GiB laptop



MemTotal:	23397872 kB
MemFree:	1394720 kB
MemAvailable:	15413048 kB
Buffers:	184188 kB
Cached:	13302188 kB
SwapCached:	41176 kB
Active:	8512216 kB
Inactive:	10476956 kB
Active(anon):	1169736 kB
Inactive(anon):	4611376 kB
Active(file):	7342480 kB
Inactive(file):	5865580 kB




# Memory metrics /proc/meminfo

- Example on my 24 GiB laptop
- Buffers: kernel buffers, for block I/O & IPC
- Cached: file pages in memory
  - ◆ Include tmpfs & shmem
  - ◆ Exclude swapcached
- Swapcached: Memory that was in swap, was put back in RAM but kept in swap

MemTotal:	22.3139	GB
MemFree:	1.35636	GB
MemAvailable:	14.693	GB
Buffers:	181.113	MB
Cached:	12.6481	GB
SwapCached:	41.8359	MB
Active:	8.1002	GB
Inactive:	9.97824	GB
Active(anon):	1.10642	GB
Inactive(anon):	4.40702	GB
Active(file):	6.99377	GB
Inactive(file):	5.57122	GB

# Memory metrics /proc/meminfo



# Memory metrics

- Mem prefixed metrics don't include swap
- Cached is a huge metric but also imprecise
  - ◆ Code in the kernel itself to check if cached < 0
  - ◆ Quite some subcomponents
  - ◆ Despite the name, everything is not "cache" memory
  - ◆ Everything can't be reclaimed
  - ◆ Can send partially to swap, but != swapcached
- Unevictable means memory that can't be sent to swap
- Mapped is mmap(2)-ed files
  - ◆ No anonymous mmap for example

# Different kinds of memory

Complex graph shown above indicates  
how complex it actually is



# Dirty memory

- When writing data to a file, by default the data isn't actually written
  - ◆ Well not directly, not always, and it's difficult to predict default behavior
- Because of performances reasons, when writing to a file, the data is actually put in a special cache in the kernel
- This cache has a special name: dirty memory
- Dirty memory is a trick played on the user:
  - ◆ We told them the data is written (`write(2)` succeeded)
  - ◆ It's actually not really on disk

# Dirty memory

- Dirty memory is “dangerous”
- A hard failure of the system, bug in the kernel, or some nasty crash, and the data it lost
- Dirty memory must be flushed down to the disk
- Dirty memory helps for performances, but introduces a risk
- In fact, MacOS and windows do this as well
  - ◆ “Don’t unplug the USB key without ejecting it”
- How to control dirty memory ?
- sync(2)


# Dirty memory & writeback

- open(2) flags like O\_DIRECT
- Check dirty memory size and watch for high or constant high values
  - ◆ It might means that disks are a bottleneck
- The cache mechanism for write is called write-back
- It works with LRU lists
  - ◆ Active and inactive list
    - To handle one-access cache eviction case
  - ◆ Known as LRU/2

# Dirty memory & writeback

- Dirty memory (and page cache in general) is implemented with a struct address\_space
- These structs are kept in a radix tree
  - ◆ Meaning that the struct are ordered in a prefix tree by their address pointer

1 romane  
2 romanus  
3 romulus  
4 rubens  
5 ruber  
6 rubicon  
7 rubicundus




# Dirty memory & writeback

- Flushing dirty pages to disk is done asynchronously
  - ◆ Unless cache is full during a cache manipulating operation
- A page is flushed when it has stayed in the cache long enough
  - ◆ Or when memory is running low
  - ◆ Or when manually requested with sync(2)
- Behavior is also tunable via knobs in /proc/sys/vm
  - ◆ There's even a laptop\_mode option !
    - Sadly mostly useless nowadays

# Dirty memory & writeback

- Flushing dirty pages to disk is done asynchronously
  - ◆ Unless cache is full during a cache manipulating operation
- A page is flushed when it has stayed in the cache long enough
  - ◆ Or when memory is running low
  - ◆ Or when manually requested with sync(2)
- Behavior is also tunable via knobs in /proc/sys/vm
  - ◆ There's even a laptop\_mode option !
    - Sadly mostly useless nowadays



# Dirty memory handling

- How to have both performances and data integrity assurance ?
- Need to trick with concepts like WAL
  - ◆ Write-ahead Logging
- Imagine a database context
- You don't want to lose data
- But transactions must be quick
  - ◆ As quick as possible
- Transactions can be complex. They can impact your whole data

# Dirty memory handling


- Likely, the database will be stored on disk on a file
- It can be huge, so a modification can introduce changes in quite some “random” places of the file
- Random access to different places of the file is expensive, in terms of I/O
- Writing the result after each transaction will take a lot of time
  - ◆ potentially

# Dirty memory handling

- Instead, WAL technique allow to deal with this exact behavior
- The WAL is a log file that will record each transactions, in the right order
- When a client makes a query:
  - ◆ The transaction is written to the WAL
  - ◆ We make sure the WAL is written to disk
  - ◆ We perform the transaction, and return the value
  - ◆ Later, the modified db file might be flushed out to disk

# Dirty memory handling

- If the database server crashes badly, the WAL is still there
- All successful transactions might not have been flushed to disk
- The database engine will check its WAL, and assure that data is correct
- If not, it can correct it since it has all the information in the WAL
- Regularly the WAL is reseted with a checkpoint
- Writing to disk the WAL is less expensive since it's append-only mode



# NUMA nodes

Architecture comes to play



# NUMA nodes

- On some architecture, not all memory is on the same access level
- Especially on “big” servers where it’s not uncommon to have 2 CPUs
  - ◆ And 2 memory zones
- Instead of having Unified Memory Access, we now have Non Unified Memory Access
- Reaching memory in node 1 from CPU on node 0 is possible
  - ◆ But more expensive

# NUMA nodes

- Linux is NUMA aware
- numactl --hardware
- cat /proc/cpuinfo ; cpuinfo
- The scheduler runs in best-effort by default
- If a task has been running in a NUMA node, it will try to keep it there
- Has some functions and data structure to perform its NUMA assignation
- <https://elixir.bootlin.com/linux/latest/source/kernel/sched/fair.c#L1439>

# NUMA nodes

- But this best-effort mode can actually be not good enough
- Especially in some cases where the machine is quite loaded
- It might actually sometimes be a good trade to force a task to run on a NUMA node
  - ◆ Reducing its CPU & RAM capacities
- Good example: VMs on an hypervisor

# NUMA nodes

- Possibility to visualize NUMA memory allocation
- numastat
  - ◆ Has even a per-process information
  - ◆ `/proc/<pid>/numa_maps`
- Possibility to set a NUMA policy
  - ◆ `set_mempolicy(2)`
  - ◆ `sched_setaffinity(2)` also
  - ◆ Or via cgroups


# Sidenote: Pushing debug tools even further

Carcinization of debugging tools



# Advanced debug tools

- strace(1) is considered a debugging tool
  - ◆ syscall oriented
- gdb(1) remains the “true and only debugger”
  - ◆ Or is it ?
- gdb(1) while having tons of functionalities lacks a critical component
- record & replay
- Mozilla introduces rr





rr

- The rr project is a debugger project built on top of GDB
- Not a replacement
  - ◆ Allow you to keep using all the GDB features
  - ◆ Not asking you to learn everything again
- rr works by recording your buggy software first
- Like strace(1), it will inspect closely what your program does
- Record it
- And provide a way to replay it, in the exact same context

- Replaying the exact same session is very useful for a few reasons:
  - ◆ No need to make the user interact the same way everytime
  - ◆ Ability to catch a misfortune once and work on it
    - Race conditions, thread problems, ...
  - ◆ Keep learned info in a debugging session across runs (pointer values, etc)

“ Let's check a quick example





rr

- rr was designed by Mozilla to debug firefox
- It's able to debug complex software like firefox
- It has some limitations though
  - ◆ Single core machine emulated
  - ◆ x86 CPU
  - ◆ some syscalls not tracked
  - ◆ Can break on kernel update

# rr - how does it work ?

- rr when started records everything to replay the exact same session
- ptrace(2)
- seccomp(2)
- Because one of the things rr tries to catch is race conditions between threads, it must be able to catch them
- rr runs all threads on the same CPU core to be sure to capture such events
  - ◆ Impacts perf on heavily parallelized computation

# rr - how does it work ?

- To be able to run all threads efficiently on the same core and catch their output, rr is preemptive
- When a thread enters a syscall, ptrace(2) catches the syscall and hand is given back to rr
- rr also periodically preempts threads with signals
- It chooses which thread to run while trying to respect linux scheduler and its priorities

# rr - how does it work ?

- When a program do a syscall, rr catches it because of `ptrace(2)`
- It chooses to resume the syscall, but catches the return value
  - ◆ Like `strace(1)`
- It stores the syscall interaction in a replayable format
- It works for most syscalls, but `ptrace(2)` itself
- A process can only be `ptrace-d` once
  - ◆ And firefox and many other already use `ptrace(2)` on themselves

# rr - how does it work ?


- rr emulates ptrace(2) syscall to bring compatibility
- rr has to deal with complex situations
  - ◆ ioctl
  - ◆ Namespaces
  - ◆ ...
- To replay a recorded trace, ptrace(2) is also used
- rr replaces all syscalls with breakpoints
  - ◆ It moves past the breakpoint, and set the return value as recorded

# rr - how does it work ?

- Some syscalls are harder to replay
  - ◆ mmap(2) - you need to have the same address
  - ◆ execve(2) - you have memory mappings that can change (ASLR)
- rr has to trick or implement complex logic to properly emulate them
- Asynchronous events must also be handled
  - ◆ Signals, interrupts
- They must be sent at the exact same time

# rr - how does it work ?

- rr is able to time precisely when async events occurs to replay them the same way
- It relies on x86 specific performance counters
- rr must also catch race conditions happening on shared memory
- As they describe, famous cases includes X server, pulseaudio, GPU related function and vdso
- They disable shared memory for X and pulseaudio and remove direct access for GPU
  - ◆ Worse perfs, but ability to replay the bug




# rr - how does it work ?

- For VDSO, rr live-patches vdso in the tracee address space to replace VDSO calls to actual syscalls
- rr must also be able to catch non-deterministic CPU instructions
- RDTSC is caught via prctl(2)
- RDRAND is rarely used, it's replaced manually in the few places found, but this is not caught by rr
- CPUID returns the core number, so sched\_setaffinity(2) is used to force a core

# rr - how does it work ?

- For VDSO, rr live-patches vdso in the tracee address space to replace VDSO calls to actual syscalls
- rr must also be able to catch non-deterministic CPU instructions
- RDTSC is caught via prctl(2)
- RDRAND is rarely used, it's replaced manually in the few places found, but this is not caught by rr
- CPUID returns the core number, so sched\_setaffinity(2) is used to force a core
  - ◆ These instructions might be handled differently in recent versions thanks to CPUID faulting



# rr - how does it work ?

- As a practical point of view, the trace created shall remain quite small
- It's compressed (and decompressed) on-the-fly by rr
- Shared libraries and binaries are stored via hard links or cow mechanisms
- Because ptrace(2) introduces a context switch (from tracee to tracer and vice-versa), and because it's used twice per syscall (before and after), it affects performances drastically
  - ◆ But rr is clever

# rr - how does it work ?

- To avoid running too many `ptrace(2)`, rr injects a library in each tracee
- The library overwrites syscall wrappers
- The library performs the syscall, but writes information to a shared buffer, shared with rr
- It tries to catch most frequently used syscalls this way
  - ◆ But fallbacks to the `ptrace(2) + syscall` in other cases
- ... there are many other challenges solved by rr
  - ◆ [Read the paper explaining most of them](#)

# rr - how does it work ?

- The master of engineering put in rr leads to a very practical tool
- The overhead it adds is about 20% on firefox
- If firefox takes 10min to perform a task, it will take 12min max with rr as observed
- All these elements make rr also very powerful with fuzzers

# Memory overcommitting

Let's go beyond limits



# Memory overcommitting

- /proc/meminfo also has some metrics about virtual memory
- On linux, you can over-allocate
  - ◆ vm.overcommit\_memory + vm.overcommit\_ratio
- An allocation in virtual memory != necessarily bound to physical memory
  - ◆ It is if it's used, meaning written to
- Useful because softwares tend to allocate more than they actually use
  - ◆ That's also a reason why you'll unlikely see a negative answer from malloc(3)

# Memory overcommitting

- 3 overcommitting modes possible:
  - ◆ 0 -> heuristic, let the kernel decide (default)
  - ◆ 1 -> always allow, never check
  - ◆ 2 -> always check
- In /proc/meminfo:
  - ◆ Committed\_AS is the sum of all committed (allocated virtual memory for all processes)
  - ◆ CommitLimit is the maximum amount of memory allocatable
    - Makes sense in mode 2 only

# Memory overcommitting

- Overcommitment lead to memory limit being hit before a memory allocation syscall fails
- Checking return value of malloc(3) won't guarantee the memory is yours
- You will trigger the OOM-killer in fact
- Still check return value ...


# memory metrics

- Other metrics are also available in:
  - ◆ /proc/vmstat
  - ◆ /proc/swap
  - ◆ /proc/buddyinfo
  - ◆ ....
- Also some per-process metrics
  - ◆ /proc/<pid>/maps
  - ◆ ...

Can we pause a  
minute and finally  
explain /proc ?


A small dive in pseudofilesystems






# What is a filesystem?

- A regular filesystem should be a well known notion
- A disk (HDD, SSD, ...) is exposed as a block device on linux
  - ◆ Special file, allows “raw” access to the disk
    - Not quite, but let’s keep this definition
- To be used as one would expect (put directories, files, etc), a filesystem must be created on the disk
- A filesystem is a data layout specs
  - ◆ A data structure
  - ◆ And its driver
  - ◆ Integrated in linux through abstraction interfaces



# What is a filesystem?

- Different kinds of filesystems with different approaches, pros and cons
  - ◆ FAT, EXT4, XFS, ZFS, BTRFS, NFS, NTFS, ...
  - ◆ Can be thought for the network (NFS, CEPHFS, GLUSTERFS, ...)
  - ◆ Can have built-in snapshot mechanisms
  - ◆ Can have a journal
  - ◆ Can support extended attributes
  - ◆ Is more or less subject to fragmentation
  - ◆ ...




# What to do with a filesystem?

- Once your disk is formatted with a filesystem, it can be used
- With windows, it's directly accessible with a letter (C:, D:, ..)
  - ◆ It's simpler for them, but also kind of stupid
  - ◆ No unified hierarchy
  - ◆ What about letter conflicts ?
- In linux, you have only one hierarchy: the Virtual FileSystem

# Linux VFS

One hierarchy to rule them all



# Linux VFS

- On linux there is no disk drive letter, only "/", the root
- Linux maintain internally the VFS, a unified file hierarchy
- You can put a disk filesystem somewhere in the VFS
  - ◆ This operation is called mounting
- Everything under the mount point will be bound to the filesystem
  - ◆ Read, writes, etc
- It's common to have the root of the VFS mounted on a disk partition
- The VFS is what you can see when "exploring files" on linux

# Linux VFS

- The VFS is the concept that allows having multiple physical storage support under the same hierarchy
- It allows an abstraction of the actual operations performed to the user




Figure 13.2 The flow of data from user-space issuing a `write()` call, through the VFS's generic system call, into the filesystem's specific write method, and finally arriving at the physical media.

# Linux VFS

- The linux VFS is tightly tied to the concept of UNIX filesystem
- It was indeed built on top of the ext2 filesystem
- A UNIX filesystem in short is built with 4 concepts:
  - ◆ Files
  - ◆ Directory entries
  - ◆ Inodes
  - ◆ Mount points

# Linux VFS

- If you want to access your USB key for example, you need:
  - ◆ To have a filesystem created on your disk, or on a partition
    - The filesystem needs to be compatible with your OS
  - ◆ To mount this filesystem somewhere in your VFS
    - If it's just to access its files, you should put it somewhere it doesn't impact your system, like /mnt

# Linux VFS – mount

- To mount a filesystem in the VFS, one can use `mount(1)`
- This command (and its underlying syscall) will take a source device, and add it in the VFS at some path
  - ◆ This means that everything that used to be on this path and below isn't directly accessible anymore
    - It is still accessible by tricking
    - Opened files stay open, and modification are propagated

# Linux VFS – mount

- The source device is usually a block device (a hard drive), but it can also be something else, like:
  - ◆ A network address, when mounting a NFS partition for example (or glusterfs, cephfs, etc)
  - ◆ A special kind of source known as a pseudo-filesystem
- You can check the supported filesystem in /proc/filesystem
  - ◆ Filesystem marked with “nodev” means that they don’t need a block device

# Pseudo filesystem

Some filesystems are not like the others



# Pseudo filesystem

- A filesystem is usually meant to store and access files
- But in Unix philosophy, everything is considered a file, even if it's not truly one
- For example, you might know the special file `/dev/zero` or `/dev/null`
- There is no such infinite file on your disk than you can read forever, or write to without it being actually written
- This is an interface the kernel exposes you

# Pseudo filesystem

- When doing a `open()` syscall, the kernel will do a few things like checking the path, permissions, etc ...
- Then it will dispatch the syscall to the driver responsible for the file
  - ◆ If the file is on an ext4 partition for example, we need to run code specific to ext4 data structure (which is in the end what a filesystem is)

# Pseudo filesystem

- We could come up with a special filesystem driver, that will execute functions for us depending on the file we read/write
- For example, a file that will execute this function when read:

```
def read_dev_zero(length, buff):  
    if len(buff) > length:  
        length = len(buff)  
    memset(buff, 0, length)  
    return length
```

This is obviously pseudo-code  
and not the actual linux  
implm of /dev/zero

# Pseudo filesystem

- We can go a bit further, and imagine this as a whole interface
- For example, /proc
- It's a pseudo filesystem mounted in /proc called procfs
- procfs exposes information about processes and various other runtime information
  - ◆ meminfo, filesystems supported, etc
- When reading a file there, you actually run kernel code that generates a response for you
- There is no disk space taken, only RAM for the responsible kernel code

# Pseudo filesystem - procfs

- procfs goal is quite easy to understand, and is mostly read-only to return kernel runtime values
- But we can have other filesystems a bit more complex
- procfs for example which role is to expose current kernel parameters and settings for many things (memory, network, etc).
  - ◆ They can be read, but also written to, to dynamically change the kernel behaviour
  - ◆ You can for example disable IPv6, drop memory caches, etc ....

# Pseudo filesystem - procfs

- procfs is more or less the config interface for the kernel, with the command line
- Regularly used with sysctl(1) binary

# Pseudo filesystem - tmpfs

- tmpfs is a very useful pseudofilesystem
- Everything inside is stored in RAM
  - ◆ Very fast accesses
  - ◆ volatile , reboot = data gone
  - ◆ Usually mounted at least in /tmp
- When mounting this pseudofilesystem, size argument used to give the maximum size
  - ◆ Defaults to half the RAM

# Pseudo filesystem - devfs

- devtmpfs is also a well known pseudo filesystem expected to be mounted on all platforms, on /dev
- It's a bit special, being a tmpfs, another pseudofilesystem, but with special behaviour
- It differs from tmpfs by having automatically registered Linux driver register block devices they create in the filesystem
- /dev – or devtmpfs – contains a block and chardevices:
  - ◆ Your disks – and their partition(s) if any
  - ◆ Special files like zero, null, urandom, kms...
  - ◆ Your tty(s)
  - ◆ ....

# Pseudo filesystem - cgroups

- Another pseudofilesystem you might have encountered already is the cgroups (v1 or v2) fs
- Interface to manipulate the control groups
- Use extensively by systemd, docker, etc ...
- Let's not get into too many details here

# Understanding VFS structure

Will be useful to understand some metrics and do advanced monitoring



# The VFS structure

- The linux VFS is tightly tied to the concept of UNIX filesystem
- It was indeed built on top of the ext2 filesystem
- A UNIX filesystem in short is built with 4 concepts:
  - ◆ Files
  - ◆ Directory entries
  - ◆ Inodes
  - ◆ Mount points

# The VFS structure

- The UNIX filesystem build those 4 representations this way:
  - ◆ File
    - A file is a set of bytes, and doesn't contain metadata
    - A directory is a special kind of file that lists its content
  - ◆ Inode
    - An inode represents the metadata of a file. It has a unique number in a given filesystem
  - ◆ Mount points - or superblock
    - Contains metadata information for the whole fs

# The VFS structure


- The UNIX filesystem build those 4 representations this way:
  - ◆ Dentry
    - Directory Entry
    - Represents the components of a path

# The VFS structure

- The VFS will be built on those 4 unix concepts
- Any filesystem not implementing a concept listed above will have to provide a compatibility layer
  - ◆ The driver will have to create one of those concept on-the-fly
  - ◆ With some (usually) negligible overhead
- Those are requirements for the VFS

# The VFS structure

- The VFS is the abstract representation exposed to the user, more or less indirectly
- It needs to be abstract and compatible with any “backend”
- Which means its structure must be able to interact with any actual filesystem implementation
  - ◆ For pseudo-filesystems implemented in linux, it's trivial, but for external ones, harder
- The VFS can be complicated with many mounts, filesystems mounted in multiple places, etc



# The VFS structure - super\_block

- The UNIX filesystem concept of a superblock is mapped to a struct super\_block in the VFS
- <https://elixir.bootlin.com/linux/latest/source/include/linux/fs.h#L1451>
- This struct contains information about a mount point
- It contains a struct super\_operations that will provide functions to filesystem-specific pointers for the filesystem-specific operations
  - ◆ <https://elixir.bootlin.com/linux/latest/source/include/linux/fs.h#L2222>

# The VFS structure - super\_block

- The super\_block is usually a mapping to the filesystem's control block of superblock, stored on the disk for regular filesystems
  - ◆ The metadata information for a filesystem
  - ◆ How many files, its size, ...
  - ◆ Generated on-the-fly for pseudo-filesystems
- Contains also run-time information for the mount-point
  - ◆ Is frozen ? Is dirty ? Mount flags, ...
    - Frozen = block write operation on a fs

# The VFS structure - super\_block

- Its operations struct will allow operations of the super\_block itself
  - ◆ Sync to the disk, remount, freeze, get statistics, ...
- But also on the inodes it handles
  - ◆ Create, delete, dirty, ...
- Having the super\_operations allows genericity in the manipulation of super\_block object, but filesystem-specific implementations of such operations

# The VFS structure - inode

- Implemented in the VFS as the struct inode
  - ◆ <https://elixir.bootlin.com/linux/latest/source/include/linux/fs.h#L593>
- Also contains a struct for an inode's operation
  - ◆ [include/linux/fs.h - Bootlin](#)
- This struct contains information about an inode (file metadata)
- The inode content is written on the disk but the struct is generated when the file is accessed

# The VFS structure - inode

- Since an inode is generic for all files, and in UNIX everything is a file, struct inode contains a union for specific files
  - ◆ i\_pipe, i\_bdev, i\_cdev
- An inode contains quite some fields that can be omitted in a driver implementation
  - ◆ For example i\_atime
- Operations includes:
  - ◆ create, mkdir, mknod, symlink, permissions, ...
- No read, write !

# The VFS structure - file

- It is important to distinguish a file from the UNIX filesystem concept to the struct file, aka the VFS file concept
- A UNIX file is what people usually understand by a file
  - ◆ Without the metadata
- A struct file represents a per-process file interaction
- A struct file is what process usually interact with
  - ◆ Especially a file.f\_op
- include/linux/fs.h - Bootlin

# The VFS structure - file

- A struct file is created by open(2)
- A process can have multiple struct files pointing to the same UNIX file, the same inode
- The struct inode is however unique
- Its operations are all the operations you think about when thinking about file manipulation from a process
  - ◆ read, write, lseek, ioctl, mmap, ...

# The VFS structure - dentry

- The difference between the kernel and users for files are how they are identified
- A user (usually) identifies a file by its path
  - ◆ open(2) syscall for example
- The kernel by an inode number
  - ◆ The path is used to translate to this inode concept
  - ◆ The same file (same inode) can have multiple paths for example
- A file is a generic term and can have multiple types, including being a directory

# The VFS structure - dentry

- Each component of a path is decomposed in objects called dentry
- /bin/bash is .. 3 dentry objects
  - ◆ /, bin and bash
  - ◆ The first 2 are dentry representing a directory, the latest is a regular file
- A dentry object is a VFS specific object. There's no direct information about underlying object pointed by it
- struct dentry

# The VFS structure - dentry

- The role of the dentry object is to ease the user manipulation of file and directories
- Those operations are costly
  - ◆ String manipulation
  - ◆ Need to check if valid
  - ◆ Check its subcomponents
  - ◆ ...
- The dentry object is really meant to represent a path
- A mount point, a directory, a file will have a struct dentry
  - ◆ When needed

# The VFS structure - dentry

- The role of the dentry object is to ease the user manipulation of file and directories
- Those operations are costly
  - ◆ String manipulation
  - ◆ Need to check if valid
  - ◆ Check its subcomponents
  - ◆ ...
- The dentry object is really meant to represent a path
- A mount point, a directory, a file will have a struct dentry
  - ◆ When needed

# The VFS structure - dentry

- A dentry can be positive or negative
  - ◆ A positive ones means it has an inode associated to it
  - ◆ A negative one is the opposite
- A negative dentry (because the path is wrong for example) can be kept in cache to resolve queries quicker
- A dentry can also be considered as used or unused via `d_count`
  - ◆ `d_count` counts the number of active reference to the associated inode
    - Meaning if there are active users of the object

# The VFS structure - dcache



- All those presented mechanisms lead to an obvious design:  
dcache
- dcache is a cache mechanism to store, access and remove  
dentry objects to have quicker accesses to files
- dcache keeps track of dentry objects, in both active (used)  
state, and inactive (unused but valid) and negative state  
(invalid)
- It provides a hash table to have quick access
  - ◆ `d_lookup()`

# The VFS structure - dentry


- The struct dentry also contains an operation struct
  - ◆ <https://elixir.bootlin.com/linux/v6.0.7/source/include/linux/dcache.h#L127>
- The operations of a dentry includes:
  - ◆ revalidate, hash, compare, ...

# struct dentry

```
  struct dentry  
  {  
    unsigned int d_flags;  
    seqcount_spinlock_t d_seq;  
    struct hlist_bl_node d_hash;  
    struct dentry *d_parent;  
    struct qstr d_name;  
    struct inode *d_inode;  
    unsigned char d_iname [DNAME_INLINE_LEN];  
    struct lockref d_lockref;  
    const struct dentry_operations *d_op;  
    struct super_block *d_sb;  
    unsigned long d_time;  
    void *d_fsdma;  
    union  
    {  
      struct list_head d_lru;  
      wait_queue_head_t *d_wait;  
    };  
    struct list_head d_child;  
    struct list_head d_subdirs;  
    union d_u  
    {  
      struct hlist_node d_alias;  
      struct hlist_bl_node d_in_lookup_hash;  
      struct rcu_head d_rcu;  
    };  
  };
```



→ All structs  
pointed by  
struct dentry





→ All structs pointed by struct dentry

- ◆ With “useless” structs removed




- All structs from struct dentry in a 2 elements range
- With "useless" structs






All structs from struct dentry in a 2 elements range

- With "useless" structs removed






For fun: All  
structs from  
struct dentry in  
a 10 elements  
range






For fun: All  
structs from  
struct dentry in  
a 50 elements  
range



# VFS observability

How can we observe what is going on  
with the VFS ?




# Observe the VFS

- The main interface is /proc/sys/fs
- We have for example dentry-state
  - ◆ Exposes the content of dentry\_stat\_t
  - ◆ <https://elixir.bootlin.com/linux/v6.0.7/source/fs/dcache.c#L118>
- file-max
  - ◆ <https://elixir.bootlin.com/linux/v6.0.7/source/include/uapi/linux/fs.h#L97>
- A few other interfaces, but yet limited


# Observe the VFS


- The VFS is hardly observable, and actually matters less than actual filesystem underneath
- Syscall to get some information about a filesystem: statfs(2)
  - ◆ Actually the recommended glibc wrapper in statvfs(3)
- Hardly any other interfaces for filesystems :/
- What is to be observed would mostly be I/O on actual physical devices
- It is frustrating because there are quite some structs, operations going on
  - ◆ But no interfaces ...



# eBPF

Let's talk about the elephant in the room





# Linux observability

- Most of this course has been focusing on two things:
  - ◆ Understanding linux kernel mechanisms
  - ◆ Understanding how to observe them
    - And observability in general
- For the past years, a fancy term appeared and is mentioned frequently
  - ◆ Everytime one talk about linux and observability in the same sentence
  - ◆ On lots of cool, modern and alpha projects

# eBPF


- eBPF is relatively new feature added to linux “recently”
- It is actually an extension of a feature existing in the kernel for years: BPF
  - ◆ Berkeley Packet Filter
    - Berkeley being the same as in Berkeley sockets or BSD
  - ◆ Now it's extended BPF

# Getting to know BPF

- eBPF is a newer and more powerful BPF
  - ◆ But what is BPF already ?
- BPF is an old project that has its roots in BSD
- It's from the early 1990
- It has been integrated in Linux in the beginning of the 2000's
  - ◆ Around linux 2.5
- It has since been evolving gradually to become what is known as eBPF

# Getting to know BPF


- BPF until the 2010's had a simple purpose: filter network paquets
- It comes with a virtual machine running in kernel land
- It executes BPF scripts
  - ◆ BPF scripts are small programs with special instructions
  - ◆ It's not x86 (or else) instructions directly
- Scripts are limited
- Used by tcpdump for example



# Getting to know BPF

● ● ●

```
1 $ sudo tcpdump -ni any dst host 127.0.0.1 and dst port 19999
2 tcpdump: data link type LINUX_SLL2
3 tcpdump: verbose output suppressed, use -v[v]... for full protocol decode
4 listening on any, link-type LINUX_SLL2 (Linux cooked v2), snapshot length 262144 bytes
5 17:42:23.798260 lo    In IP 127.0.0.1.46842 > 127.0.0.1.19999: Flags [S], seq 3512779186, win 65495, options
   [mss 65495,sackOK,TS val 2009782134 ecr 0,nop,wscale 7], length 0
6 17:42:23.798328 lo    In IP 127.0.0.1.46842 > 127.0.0.1.19999: Flags [.], ack 908325355, win 512, options
   [nop,nop,TS val 2009782134 ecr 2009782134], length 0
7 17:42:23.798508 lo    In IP 127.0.0.1.46842 > 127.0.0.1.19999: Flags [P.], seq 0:79, ack 1, win 512, options
   [nop,nop,TS val 2009782134 ecr 2009782134], length 79
8 17:42:23.798651 lo    In IP 127.0.0.1.46842 > 127.0.0.1.19999: Flags [.], ack 32769, win 379, options
   [nop,nop,TS val 2009782135 ecr 2009782135], length 0
9 17:42:23.799084 lo    In IP 127.0.0.1.46842 > 127.0.0.1.19999: Flags [F.], seq 79, ack 40380, win 512, options
   [nop,nop,TS val 2009782135 ecr 2009782135], length 0
10 ^C
11 5 packets captured
12 10 packets received by filter
13 0 packets dropped by kernel
```




# Getting to know BPF



```
1 $ sudo tcpdump -ni any dst host 127.0.0.1 and dst port 19999 -d
2 tcpdump: data link type LINUX_SLL2
3 (000) ldh      [0]
4 (001) jeq      #0x800          jt 2    jf 14
5 (002) ld      [36]
6 (003) jeq      #0x7f000001    jt 4    jf 14
7 (004) ldb      [29]
8 (005) jeq      #0x84           jt 8    jf 6
9 (006) jeq      #0x6            jt 8    jf 7
10 (007) jeq     #0x11           jt 8    jf 14
11 (008) ldh     [26]
12 (009) jset    #0x1fff         jt 14   jf 10
13 (010) ldxb    4*([20]&0xf)
14 (011) ldh     [x + 22]
15 (012) jeq     #0x4e1f          jt 13   jf 14
16 (013) ret     #262144
17 (014) ret     #0
```

# Getting to know BPF

- Executing code dynamically in the kernel in a virtual machine looks very interesting
- One could want to execute more than just packet filtering
- Maybe report information about the packets to userland ?
- Maybe include some packet manipulation ?
  - ◆ Even dropping them ?
- Etc ...
- As ideas came along, BPF subsystem grew



# From BPF to eBPF: internal BPF

- The virtual machine idea looked seducing
- A JIT was introduced to compile BPF instruction to x86 instructions
  - ◆ So it was fast
- BPF had too many limitations, like the number of registers available (2), no 64 bits registers, ...
- 2 BPFs were created: classic BPF and internal BPF
- Internal BPF was more powerful, but was hidden from userspace

# From BPF to eBPF: internal BPF

- The only interface was for classic BPF
- Classic BPF was then transformed in internal BPF in the kernel
  - ◆ Faster, using x86 instructions directly, ...
- But the idea of exposing an interface for internal BPF to userland was already there
- The goal was also to work on GCC/LLVM to generate internal BPF directly
- Internal BPF was also able to call some limited set of kernel functions

# From BPF to eBPF: internal BPF

- Internal BPF was also starting to draw attention for tracing
- Could be useful to use this dynamic language and virtual machine to run more things dynamically
  - ◆ No more real connection with the network
- BPF moved from the net/ to kernel/bpf
  - ◆ Removed also some ties with network as well
- This internal BPF was then renamed as eBPF for extended BPF

# eBPF evolution

- What happened after ?
- bpf(2) syscall was introduced
- Load a BPF program in the kernel
- Of course, huge restrictions:
  - ◆ We can't run arbitrary code in the kernel
  - ◆ We can't call any function
  - ◆ We can't have a loop
  - ◆ We can't sleep

# eBPF evolution

- The introduction of the bpf(2) syscall also introduce a BPF verifier:
  - ◆ Tries to prevent harmful programs from being loaded
  - ◆ Things mentioned on the previous slide
  - ◆ But also read from unallocated registers, bound checks, etc
- Obviously a privileged syscall requiring CAP\_SYS\_ADMIN (until 5.8 when CAP\_BPF was added)

# eBPF evolution

- Added feature to eBPF at that time was also maps
- eBPF maps are key/value store to exchange data from kernel to userland
  - ◆ They are created from userland though
    - Via a call to bpf(2) too
    - But not directly from a running eBPF program

# eBPF evolution

- eBPF evolutions since 2014
  - ◆ Adding persistent eBPF programs
  - ◆ Adding a pseudo filesystem for eBPF (called bpf)
  - ◆ Adding different types of eBPF programs
    - Few related to network traffic
    - One to change socket(2) types
    - Few related to monitoring
  - ◆ Adding types of maps

# eBPF evolution since 2014

- Adding support to authz sysctl via eBPF
- Ability to dump kernel structures
- Allow sleepable eBPF programs
- Ability to call some restricted kernel functions
- Control scheduler decisions
- Allow to loop (with still quite some restrictions)
- ....

# Running eBPF

How can one create and run such  
magical programs ?



# How to run eBPF ?

- An eBPF program is setup in the kernel via bpf(2)
- It is checked, verified and then installed
- It can be referenced via an ID
- But a call to bpf(2) doesn't run the program
- An eBPF program doesn't run when userland asks for it to run
- An eBPF program is linked to an event, and is started from this event
  - ◆ Like classic BPF, when a socket receives data for example


# The bpf(2) syscall

- All interactions with (e)bpf goes through the bpf(2) syscall
- The design intended to have only a single syscall for all operations
  - ◆ Like ioctl(2)
- First argument is the cmd:
  - ◆ Create, read elem, update elem, delete elem for maps
  - ◆ Load a eBPF program

# The bpf(2) syscall

- A map to create must also have a type
  - ◆ BPF\_MAP\_TYPE\_HASH
  - ◆ BPF\_MAP\_TYPE\_ARRAY
  - ◆ BPF\_MAP\_TYPE\_STACK\_TRACE
  - ◆ ...
- All map types don't behave the same way
- Read the documentation for specifics
- Same goes for eBPF programs
  - ◆ Different types for different usages
  - ◆ Changes in the verifier and capacities of your program

“ *Ok but why the fuss  
around eBPF ?*



# eBPF in action

Checkout why eBPF rhymes with observability



# Linux kernel's own javascript

- eBPF is great because it adds dynamicity to the kernel
  - ◆ Could be seen as the equivalent of javascript, but for the kernel
- There's no intention of putting cats animation in the kernel though
- But instrumentality and observability is another subject
- Let's discover a few eBPF project


# eBPF projects

- Netdata kernel collector
  - ◆ Collects metrics and allow monitoring on events that were inaccessible so far
  - ◆ Process-related, VFS, hardirqs, softirqs, shmem, sync-related syscalls, file access, mount, network-related, TCP-related internal functions calls, ...

# eBPF projects

- iovisor/bcc
  - ◆ Toolkit to manipulate eBPF easily: write eBPF programs in C-like language and compiled with LLVM, front-end for eBPF programs with python or lua
  - ◆ Comes with pre-defined tools to monitor, trace, snoop a machine

# Linux bcc/BPF Tracing Tools



# eBPF projects


- Cilium
  - ◆ Kubernetes-related projects for network
  - ◆ CNI to bring eBPF-aware networking to Kubernetes with:
    - Loadbalancing
    - Network policy L7 aware
    - ...
  - ◆ Hubble for observability in Kubernetes networking related stack
    - Metrics, tracing

# eBPF projects

- Check a bigger, updated and more detailed list on  
<https://ebpf.io/applications/>
- Those projects look awesome and very promising
- The ability to expose metrics un-exposable otherwise is astonishing
- But how exactly is this possible ?
  - ◆ Let's stop with the vagueness around observability and eBPF and let's dig into implementation details

# kprobes, uprobes, tracepoints, ...

They were there all along !




# Probes and tracepoints

- Before “modern” monitoring like eBPF allow us to do, there were already concepts in the kernel to get events
  - ◆ From the early 2000’s
- It was mostly for instrumentation and debugging than observability
- Mostly aimed for kernel developers at first
- Then brought to more people


# Probes and tracepoints

- There are 2 categories of event source in the kernel
  - ◆ Dynamically defined tracing points (Probes)
  - ◆ Statically defined tracing points (Tracepoints)
- Linux offers:
  - ◆ Tracepoints
  - ◆ Kprobes
- But also for userland:
  - ◆ Uprobes
  - ◆ (USDT)

# Probes and tracepoints



# Probes and tracepoints



# Probes and tracepoints

- Probes and Tracepoints are to collect “data”
- The way they are used, called and how the data is then exposed depends on the tracing framework
- eBPF is, among other things, a tracing framework
- iovisor/bcc presented briefly is a front-end

# Probes and tracepoints

- What is the difference between kprobes and tracepoints ?
- Tracepoints are defined statically
  - ◆ TRACE\_EVENT macro in the kernel
- They had no overhead if disabled
  - ◆ Except for a small comparison
- Once enabled, notify with info observers

# Probes and tracepoints

- What is the difference between kprobes and tracepoints ?
- kprobes are defined dynamically
- They don't require a "kprobe" event to be defined in the code
- You can compare it to a breakpoint with your debugger
- You can place it almost everywhere
  - ◆ Beginning and end of functions via k(ret)probe
- It replaces an instruction to be executed by an INT3
- The kprobe handler will check from where the trap comes from
- It will then report what is needed to the kprobe subscriber(s)

# Tracing framework

A probe on its own is hardly usable




# Tracing framework

- Probes and tracepoints reports data to a subscriber
- The subscriber is defined by the tracing framework used
- One example: eBPF
- eBPF allow you to define a small program and attach it to a probe/tracepoint
- Once the probe is fired, it calls your eBPF program
  - ◆ Argument to the probes are forwarded to your program
  - ◆ You can sometimes instrument them
  - ◆ Or do some logic and report things


# Tracing framework

- Only a limited amount of tracing frameworks are available in linux
- eBPF, ftrace and perf\_event are the 3 main choices
  - ◆ There are also out-of-tree options (SystemTap, Ittng, ...)


# ftrace




# perf\_event




# eBPF



## BPF For Tracing



Adapted from Brendan Gregg's blog post - "Linux Extended BPF (eBPF) Tracing Tools"



# Frontends

Let's build some fancy tool on top of  
these




# Frontends

- There are multiple frontends options for each framework
- The frontend usually is meant to leverage the framework easily
- Write human-readable code, and compile it in eBPF bytecode for example
- Provide awk-like scripts
- Essentially a simplification of the interfaces and syscalls
  - ◆ Sometimes shipped with a library in a given language

# Frontends

- Quite some frontends can be named:
  - ◆ Perf (for perf framework and ftrace)
  - ◆ trace-cmd
  - ◆ Bcc
  - ◆ Bpftrace
  - ◆ LTTng
  - ◆ SystemTap
  - ◆



# nsswitch digression

Understanding glibc behavior as seen by  
strace(1)



# nsswitch

- Unix world offers a few files to handle its system configuration
- Examples of those includes /etc/passwd, /etc/group, /etc/hosts, ...
- While those files works great and suit basic behavior, there are still a bit limited
- What if we wanted to handle servers' access for the employees of a company ?
  - ◆ There are hundreds of employees, thousands of systems
  - ◆ Handling each system individually is difficult and tedious

# nsswitch

- How can we extend this behavior to use other kind of services in order to provide those information ?
- For example, connect to a database to get user information
  - ◆ LDAP is a famous protocol for this
- How to handle a DNS system a bit more clever than a simple /etc/hosts + /etc/resolv.conf ?
  - ◆ With cache
  - ◆ With per-interface domain resolution for example
  - ◆ ...

# nsswitch

- GNU C Library allow us to extend and change the default behavior via a configuration file, /etc/nsswitch.conf
- nsswitch, for Name Service Switch in part of the glibc
  - ◆ And also introduced in other software due to its popularity
- The /etc/nsswitch.conf allow to change the configuration on how to find such Name Service information
- It has a pluggable approach, with shared libraries
  - ◆ Anyone can write a plugin to plug in the nsswitch system

# nsswitch.conf(5)

- Default /etc/nsswitch.conf contains the basic configuration to use the default plugins for traditional UNIX config files

# nsswitch.conf(5)

```
1 $ cat /etc/nsswitch.conf
2 # /etc/nsswitch.conf
3 #
4 # Example configuration of GNU Name Service Switch functionality.
5 # If you have the `glibc-doc-reference' and `info' packages installed, try:
6 # `info libc "Name Service Switch"' for information about this file.
7
8 passwd:      files
9 group:       files
10 shadow:     files
11 gshadow:    files
12
13 hosts:      files dns
14 networks:   files
15
16 protocols:  db files
17 services:   db files
18 ethers:     db files
19 rpc:        db files
20
21 netgroup:   nis
```

# nsswitch.conf(5)

- Default /etc/nsswitch.conf contains the basic configuration to use the default plugins for traditional UNIX config files
- It has a simple format:
  - ◆ Name service: <plugin 1> <plugin 2> ...
    - There are some limited option to add on each plugin also
- Let's check a few classic configurations

# nsswitch.conf(5)

- passwd: files systemd ldap
- For the passwd name service, first is to check with the files plugin
- The files plugin is implemented via /usr/lib/libnss\_files.so.2
- It implements the default UNIX behavior, by looking in /etc/passwd
- The next data source is implemented by libnss\_systemd.so.2
  - ◆ It implements a connector to ask systemd(1) or some specific systemd service via a systemd API information
  - ◆ nss-systemd(8)

# nsswitch.conf(5)

- passwd: files systemd ldap
- ldap is provided by nslcd and its libnss\_ldap.so.2
- Used to query nslcd daemon which connect to remote configurable LDAP server and gets users, passwords and groups from
- Because of the multiple ways of finding passwd information (different name services), cat /etc/passwd is not enough
- Prefer using getent passwd

# nscd digression

Understanding glibc behavior as seen by  
strace(1)



# nsqd


- On top of described behavior by nsswitch and glibc, another mechanism exists to provide cache for name service queries
- While /etc/passwd file for example is pretty much inexpensive to read, DNS queries or LDAP connection are expensive
- Having cache for them is great
- It's the role of nsqd to provide such cache
  - ◆ Hence its name, Name Service Cache Daemon

# nsqd

- nsqd as its name indicates is a daemon
  - ◆ It might not be installed on your machine, or not running
- It exposes a UNIX socket in /var/run/nsqd/socket
- By default, the glibc connects to this socket automatically
  - ◆ Before contacting a name service source as provided by /etc/nsswitch.conf
- If the socket can't be opened, it ... retries a second time
- If nsqd is not running, or doesn't have the info in cache, it falls back to the default nsswitch mechanism


# nscd

```
 1 socket(AF_UNIX, SOCK_STREAM|SOCK_CLOEXEC|SOCK_NONBLOCK, 0) = 3
 2 connect(3, {sa_family=AF_UNIX, sun_path="/var/run/nscd/socket"}, 110) = 0
 3 sendto(3, "\2\0\0\0\v\0\0\0\7\0\0\0passwd\0", 19, MSG_NOSIGNAL, NULL, 0) = 19
 4 poll([{fd=3, events=POLLIN|POLLERR|POLLHUP}], 1, 5000) = 1 ({fd=3,
    revents=POLLIN|POLLHUP}})
 5 recvmsg(3, {msg_name=NULL, msg_namelen=0, msg iov=[{iov_base="passwd\0", iov_len=7},
    {iov_base="\3100\3\0\0\0\0\0", iov_len=8}], msg iovlen=2, msg_control=[{cmsg_len=20,
      cmsg_level=SOL_SOCKET, cmsg_type=SCM_RIGHTS, cmsg_data=[4]}], msg_controllen=20,
    msg_flags=MSG_CMSG_CLOEXEC}, MSG_CMSG_CLOEXEC) = 15
 6 mmap(NULL, 217032, PROT_READ, MAP_SHARED, 4, 0) = 0x7fae19b4b000
 7 getrandom("\x76\x85\x0c\xee\x32\xae\x07\x34", 8, GRND_NONBLOCK) = 8
 8 brk(NULL) = 0x55ae3363c000
 9 brk(0x55ae3365d000) = 0x55ae3365d000
10 close(4) = 0
11 close(3) = 0
```



# nscd

- In previous example
- nscd is answering
- Answers with a pointer to a shared memory to mount, that contains the asked database
- mmap(2) right under





```
● ● ●
```

```
1 socket(AF_UNIX, SOCK_STREAM|SOCK_CLOEXEC|SOCK_NONBLOCK, 0) = 3
2 connect(3, {sa_family=AF_UNIX, sun_path="/var/run/nscd/socket"}, 110) = -1 ENOENT (No such file or directory)
3 close(3) = 0
4 socket(AF_UNIX, SOCK_STREAM|SOCK_CLOEXEC|SOCK_NONBLOCK, 0) = 3
5 connect(3, {sa_family=AF_UNIX, sun_path="/var/run/nscd/socket"}, 110) = -1 ENOENT (No such file or directory)
6 close(3) = 0
7 getrandom("\x0e\x3e\xf0\xab\x82\xb2\xc4\x99", 8, GRND_NONBLOCK) = 8
8 brk(NULL) = 0x55c619534000
9 brk(0x55c619555000) = 0x55c619555000
10 newfstatat(AT_FDCWD, "/etc/nsswitch.conf", {st_mode=S_IFREG|0644, st_size=359, ...}, 0) = 0
11 newfstatat(AT_FDCWD, "/", {st_mode=S_IFDIR|0755, st_size=4096, ...}, 0) = 0
12 openat(AT_FDCWD, "/etc/nsswitch.conf", O_RDONLY|O_CLOEXEC) = 3
13 newfstatat(3, "", {st_mode=S_IFREG|0644, st_size=359, ...}, AT_EMPTY_PATH) = 0
14 read(3, "# Name Service Switch configurat..., 4096) = 359
15 read(3, "", 4096) = 0
16 newfstatat(3, "", {st_mode=S_IFREG|0644, st_size=359, ...}, AT_EMPTY_PATH) = 0
17 close(3) = 0
18 openat(AT_FDCWD, "/etc/passwd", O_RDONLY|O_CLOEXEC) = 3
19 newfstatat(3, "", {st_mode=S_IFREG|0644, st_size=1800, ...}, AT_EMPTY_PATH) = 0
20 lseek(3, 0, SEEK_SET) = 0
21 read(3, "root:x:0:/:/root:/usr/bin/zsh\nb...", 4096) = 1800
22 close(3) = 0
```

# Thanks !

Questions ?





Slides available on [zarak.fr/](http://zarak.fr/)

Contact: [cyril@cri.epita.fr](mailto:cyril@cri.epita.fr)

zarak production#5492

# Sources

[lwn.net/Articles/904892/](https://lwn.net/Articles/904892/) - The ABI status of ELF hash tables [LWN.net]

[lwn.net/Articles/330589/](https://lwn.net/Articles/330589/) - KSM tries again [LWN.net]

[lwn.net/Articles/229096/](https://lwn.net/Articles/229096/) - SLUB: The unqueued slab allocator V6 [LWN.net]

[lwn.net/Articles/288056/](https://lwn.net/Articles/288056/) - TASK\_KILLABLE [LWN.net]

[lwn.net/Articles/379748/](https://lwn.net/Articles/379748/) - Huge pages part 5: A deeper look at TLBs and costs [LWN.net]

[lwn.net/Articles/317814/](https://lwn.net/Articles/317814/) - Taming the OOM killer [LWN.net]

[lwn.net/Articles/83588/](https://lwn.net/Articles/83588/) - 2.6 swapping behavior [LWN.net]

[lwn.net/Articles/82759/](https://lwn.net/Articles/82759/) - Kernel development [LWN.net]

[lwn.net/Articles/793073/](https://lwn.net/Articles/793073/) - (hopefully) saner refcounting for mountpoint dentries [LWN.net]

[lwn.net/Articles/330985/](https://lwn.net/Articles/330985/) - driver-core: devtmpfs - driver core maintained /dev tmpfs [LWN.net]

[lwn.net/Articles/331818/](https://lwn.net/Articles/331818/) - The return of devfs [LWN.net]

[lwn.net/Articles/612878/](https://lwn.net/Articles/612878/) - The BPF system call API, version 14 [LWN.net]

# Sources

[lwn.net/Articles/870269/](https://lwn.net/Articles/870269/) - Taming the BPF superpowers [LWN.net]

[lwn.net/Articles/664688/](https://lwn.net/Articles/664688/) - Persistent BPF objects [LWN.net]

[lwn.net/Articles/740157/](https://lwn.net/Articles/740157/) - A thorough introduction to eBPF [LWN.net]

[lwn.net/Articles/779120/](https://lwn.net/Articles/779120/) - Concurrency management in BPF [LWN.net]

[lwn.net/Articles/787856/](https://lwn.net/Articles/787856/) - BPF: what's good, what's coming, and what's needed [LWN.net]

[lwn.net/Articles/785263/](https://lwn.net/Articles/785263/) - Managing sysctl knobs with BPF [LWN.net]

[lwn.net/Articles/803890/](https://lwn.net/Articles/803890/) - Filesystem sandboxing with eBPF [LWN.net]

[lwn.net/Articles/818714/](https://lwn.net/Articles/818714/) - Dumping kernel data structures with BPF [LWN.net]

[lwn.net/Articles/825415/](https://lwn.net/Articles/825415/) - Sleepable BPF programs [LWN.net]

[lwn.net/Articles/856005/](https://lwn.net/Articles/856005/) - Calling kernel functions from BPF [LWN.net]

[lwn.net/Articles/873244/](https://lwn.net/Articles/873244/) - Controlling the CPU scheduler with BPF [LWN.net]

[lwn.net/Articles/794934/](https://lwn.net/Articles/794934/) - Bounded loops in BPF for the 5.3 kernel [LWN.net]

[lwn.net/Articles/132196/](https://lwn.net/Articles/132196/) - An introduction to KProbes [LWN.net]

# Sources

[lwn.net/Articles/346470/](https://lwn.net/Articles/346470/) - Fun with tracepoints [LWN.net]

[lwn.net/Articles/379903/](https://lwn.net/Articles/379903/) - Using the TRACE\_EVENT() macro (Part 1) [LWN.net]

[github.com/brendangregg/perf-tools/blob/master/kernel/kprobe](https://github.com/brendangregg/perf-tools/blob/master/kernel/kprobe) - perf-tools/kprobe at master · brendangregg/perf-tools

[jvns.ca/blog/2017/07/05/linux-tracing-systems/](https://jvns.ca/blog/2017/07/05/linux-tracing-systems/) - Linux tracing systems & how they fit together

[terenceli.github.io/%E6%8A%80%E6%9C%AF/2020/08/05/tracing-basic](https://terenceli.github.io/%E6%8A%80%E6%9C%AF/2020/08/05/tracing-basic) - Linux tracing - kprobe, uprobe and tracepoint

[github.com/iovisor/bcc](https://github.com/iovisor/bcc) - iovisor/bcc: BCC - Tools for BPF-based Linux IO analysis, networking, monitoring, and more

[oreilly.com/library/view/understanding-the-linux/0596005652/ch04s07.html](https://oreilly.com/library/view/understanding-the-linux/0596005652/ch04s07.html) - 4.7. Softirqs and Tasklets - Understanding the Linux Kernel, 3rd Edition [Book]

[docs.kernel.org/scheduler/index.html](https://docs.kernel.org/scheduler/index.html) - Linux Scheduler — The Linux Kernel documentation

[github.com/0xAX/linux-insides](https://github.com/0xAX/linux-insides) - GitHub - 0xAX/linux-insides: A little bit about a linux kernel

[kernel.org/doc/html/latest/admin-guide/mm/hugetlbpage.html](https://kernel.org/doc/html/latest/admin-guide/mm/hugetlbpage.html) - HugeTLB Pages — The Linux Kernel documentation

[kernel.org/doc/html/latest/admin-guide/mm/idle\\_page\\_tracking.html](https://kernel.org/doc/html/latest/admin-guide/mm/idle_page_tracking.html) - Idle Page Tracking — The Linux Kernel documentation

# Sources

[baeldung.com/linux/process-states](http://baeldung.com/linux/process-states) - Linux Process States | Baeldung on Linux

[haydenjames.io/what-is-iowait-and-linux-performance/](http://haydenjames.io/what-is-iowait-and-linux-performance/) - What is iowait and how does it affect Linux performance?

[unix.stackexchange.com/questions/16738/when-a-process-will-go-to-d-state](http://unix.stackexchange.com/questions/16738/when-a-process-will-go-to-d-state) - linux - When a process will go to 'D' state? - Unix & Linux Stack Exchange

[stackoverflow.com/questions/71862781/how-to-make-a-process-to-enter-d-state](http://stackoverflow.com/questions/71862781/how-to-make-a-process-to-enter-d-state) - linux - How to make a process to enter D state? - Stack Overflow

[unix.stackexchange.com/questions/539733/what-are-the-non-numeric-irqs-in-proc-interrupts](http://unix.stackexchange.com/questions/539733/what-are-the-non-numeric-irqs-in-proc-interrupts) - linux kernel - What are the non-numeric IRQs in /proc/interrupts? - Unix & Linux Stack Exchange

[elixir.bootlin.com/linux/latest/source/arch/x86/include/asm/irq\\_vectors.h](http://elixir.bootlin.com/linux/latest/source/arch/x86/include/asm/irq_vectors.h) - irq\_vectors.h - arch/x86/include/asm/irq\_vectors.h - Linux source code (v5.19.2) - Bootlin

[blog.dixo.net/irq.png](http://blog.dixo.net/irq.png)

[alexonlinux.com/smp-affinity-and-proper-interrupt-handling-in-linux](http://alexonlinux.com/smp-affinity-and-proper-interrupt-handling-in-linux) - SMP affinity and proper interrupt handling in Linux - Alex on Linux

[kernel.org/doc/html/latest/admin-guide/mm/ksm.html](http://kernel.org/doc/html/latest/admin-guide/mm/ksm.html) - Kernel Samepage Merging — The Linux Kernel documentation

# Sources

[kernel.org/doc/html/latest/RCU/whatisRCU.html](http://kernel.org/doc/html/latest/RCU/whatisRCU.html) - What is RCU? – “Read, Copy, Update” — The Linux Kernel documentation

[zeph1912.github.io/notes\\_and\\_journal\\_repo/kernel\\_softirq.html](http://zeph1912.github.io/notes_and_journal_repo/kernel_softirq.html) - Softirq | Zephyr's study notes

[unix.stackexchange.com/questions/591243/counting-the-number-of-issued-syscalls](http://unix.stackexchange.com/questions/591243/counting-the-number-of-issued-syscalls) -

[linux - Counting the number of issued syscalls](http://linux - Counting the number of issued syscalls) - Unix & Linux Stack Exchange

[rr-project.org/](http://rr-project.org/) - rr: lightweight recording & deterministic debugging

[man7.org/linux/man-pages/man7/shm\\_overview.7.html](http://man7.org/linux/man-pages/man7/shm_overview.7.html) - shm\_overview(7) - Linux manual page

[Linux-kvm.org](http://Linux-kvm.org) - KVM

[man7.org/linux/man-pages/man2/prctl.2.html](http://man7.org/linux/man-pages/man2/prctl.2.html) - prctl(2) - Linux manual page

[man7.org/linux/man-pages/man5/core.5.html](http://man7.org/linux/man-pages/man5/core.5.html) - core(5) - Linux manual page

[kernel.org/doc/html/latest/admin-guide/mm/index.html](http://kernel.org/doc/html/latest/admin-guide/mm/index.html) - Memory Management — The Linux Kernel documentation

[kernel.org/doc/html/latest/admin-guide/mm/concepts.html](http://kernel.org/doc/html/latest/admin-guide/mm/concepts.html) - Concepts overview — The Linux Kernel documentation

[lwn.net/Articles/306704/](http://lwn.net/Articles/306704/) - /dev/ksm: dynamic memory sharing [LWN.net]

[kernel.org/doc/html/latest/admin-guide/mm/numa\\_memory\\_policy.html](http://kernel.org/doc/html/latest/admin-guide/mm/numa_memory_policy.html) - NUMA Memory Policy — The Linux Kernel documentation

# Sources

[kernel.org/doc/html/latest/admin-guide/mm/numaperf.html](http://kernel.org/doc/html/latest/admin-guide/mm/numaperf.html) - NUMA Locality — The Linux Kernel documentation

[kernel.org/doc/html/latest/admin-guide/mm/pagemap.html](http://kernel.org/doc/html/latest/admin-guide/mm/pagemap.html) - Examining Process Page Tables — The Linux Kernel documentation

[kernel.org/doc/html/latest/admin-guide/mm/swap numa.html](http://kernel.org/doc/html/latest/admin-guide/mm/swap numa.html) - Automatically bind swap device to numa node — The Linux Kernel documentation

[kernel.org/doc/html/latest/admin-guide/mm/zswap.html](http://kernel.org/doc/html/latest/admin-guide/mm/zswap.html) - zswap — The Linux Kernel documentation

[stack overflow .com/questions/7880784/what-is-rss-and-vsZ-in-linux-memory-management](http://stackoverflow.com/questions/7880784/what-is-rss-and-vsZ-in-linux-memory-management) - What is RSS and VSZ in Linux memory management - Stack Overflow

[tothenew.com/blog/understanding-memory-utilization-in-linux/](http://tothenew.com/blog/understanding-memory-utilization-in-linux/) - Understanding Memory Utilization in Linux | TO THE NEW Blog

[man7.org/linux/man-pages/man7/user\\_namespaces.7.html](http://man7.org/linux/man-pages/man7/user_namespaces.7.html) - user\_namespaces(7) - Linux manual page

[redhat.com/sysadmin/dissecting-free-command](http://redhat.com/sysadmin/dissecting-free-command) - Dissecting the free command: What the Linux sysadmin needs to know | Enable Sysadmin

[kernel.org/doc/gorman/html/understand/understand011.html](http://kernel.org/doc/gorman/html/understand/understand011.html) - Slab Allocator

[man7.org/linux/man-pages/man5/slabinfo.5.html](http://man7.org/linux/man-pages/man5/slabinfo.5.html) - slabinfo(5) - Linux manual page

# Sources

[linuxize.com/post/free-command-in-linux/](http://linuxize.com/post/free-command-in-linux/) - Free Command in Linux | Linuxize

[man7.org/linux/man-pages/man5/procfs.5.html](http://man7.org/linux/man-pages/man5/procfs.5.html) - proc(5) - Linux manual page

[tecmint.com/linux-process-management/](http://tecmint.com/linux-process-management/) - All You Need To Know About Processes in Linux [Comprehensive Guide]

[devm.io/programming/linux-process-states-173858](http://devm.io/programming/linux-process-states-173858) - What are the process states in Unix/Linux?

[stackoverflow.com/questions/67769737/check-if-the-process-in-in-running-state-or-runnable-state-in-linux](http://stackoverflow.com/questions/67769737/check-if-the-process-in-in-running-state-or-runnable-state-in-linux) - Check if the process in in RUNNING state or RUNNABLE state in Linux - Stack Overflow

[elixir.bootlin.com/linux/latest/ident/task\\_struct](http://elixir.bootlin.com/linux/latest/ident/task_struct) - task\_struct identifier - Linux source code (v5.19.5) - Bootlin

[elixir.bootlin.com/linux/latest/source/include/linux/sched.h](http://elixir.bootlin.com/linux/latest/source/include/linux/sched.h) - sched.h - include/linux/sched.h - Linux source code (v5.19.5) - Bootlin

[stackoverflow.com/questions/22101574/how-to-figure-out-if-process-is-really-running-or-waiting-to-run-on-linux](http://stackoverflow.com/questions/22101574/how-to-figure-out-if-process-is-really-running-or-waiting-to-run-on-linux) - how to figure out if process is really running or waiting to run on Linux? - Stack Overflow

[eklitzke.org/uninterruptible-sleep](http://eklitzke.org/uninterruptible-sleep) - Uninterruptible Sleep

[elixir.bootlin.com/linux/latest/A/ident/TASK\\_UNINTERRUPTIBLE](http://elixir.bootlin.com/linux/latest/A/ident/TASK_UNINTERRUPTIBLE) - TASK\_UNINTERRUPTIBLE identifier - Linux source code (v5.19.5) - Bootlin

[opensource.com/article/19/2/fair-scheduling-linux](http://opensource.com/article/19/2/fair-scheduling-linux) - CFS: Completely fair process scheduling in Linux

# Sources

[docs.kernel.org/scheduler/sched-design-CFS.html](http://docs.kernel.org/scheduler/sched-design-CFS.html) - CFS Scheduler — The Linux Kernel documentation

[Oxax.gitbooks.io/linux-insides/content/Interrupts/linux-interrupts-9.html](http://Oxax.gitbooks.io/linux-insides/content/Interrupts/linux-interrupts-9.html) - Softirq, Tasklets and Workqueues · Linux Inside

[sites.google.com/site/masumzh/articles/x86-architecture-basics/interrupts-faults-and-traps](http://sites.google.com/site/masumzh/articles/x86-architecture-basics/interrupts-faults-and-traps) - Masum Z Hasan, PhD - X86 Architecture basics: Interrupts, Faults and Traps and IO

[wiki.osdev.org/Interrupts](http://wiki.osdev.org/Interrupts) - Interrupts - OSDev Wiki

[en.wikibooks.org/wiki/X86\\_Assembly/Programmable\\_Interrupt\\_Controller](http://en.wikibooks.org/wiki/X86_Assembly/Programmable_Interrupt_Controller) - x86 Assembly/Programmable Interrupt Controller - Wikibooks, open books for an open world

[intel.com/content/dam/www/public/us/en/documents/manuals/64-ia-32-architectures-software-developer-vol-3a-part-1-manual.pdf](http://intel.com/content/dam/www/public/us/en/documents/manuals/64-ia-32-architectures-software-developer-vol-3a-part-1-manual.pdf) - 64-ia-32-architectures-software-developer-vol-3a-part-1-manual.pdf

[man7.org/linux/man-pages/man2/madvise.2.html](http://man7.org/linux/man-pages/man2/madvise.2.html) - madvise(2) - Linux manual page

[kernel.org/doc/html/latest/admin-guide/mm/transhuge.html](http://kernel.org/doc/html/latest/admin-guide/mm/transhuge.html) - Transparent Hugepage Support — The Linux Kernel documentation

[docs.kernel.org/admin-guide/mm/index.html](http://docs.kernel.org/admin-guide/mm/index.html) - Memory Management — The Linux Kernel documentation

[kernel.org/doc/html/latest/filesystems/vfs.html](http://kernel.org/doc/html/latest/filesystems/vfs.html) - Overview of the Linux Virtual File System — The Linux Kernel documentation

# Sources

[books.gigatux.nl/mirror/kerneldevelopment/0672327201/ch12lev1sec7.html](http://books.gigatux.nl/mirror/kerneldevelopment/0672327201/ch12lev1sec7.html) - The Dentry Object  
[developer.ibm.com/tutorials/l-completely-fair-scheduler/](http://developer.ibm.com/tutorials/l-completely-fair-scheduler/) - Inside the Linux 2.6 Completely Fair Scheduler